

DIRECT VOLUME RENDERING

PART I: FOUNDATIONS, TRANSFER FUNCTIONS, LIGHTING MODELS

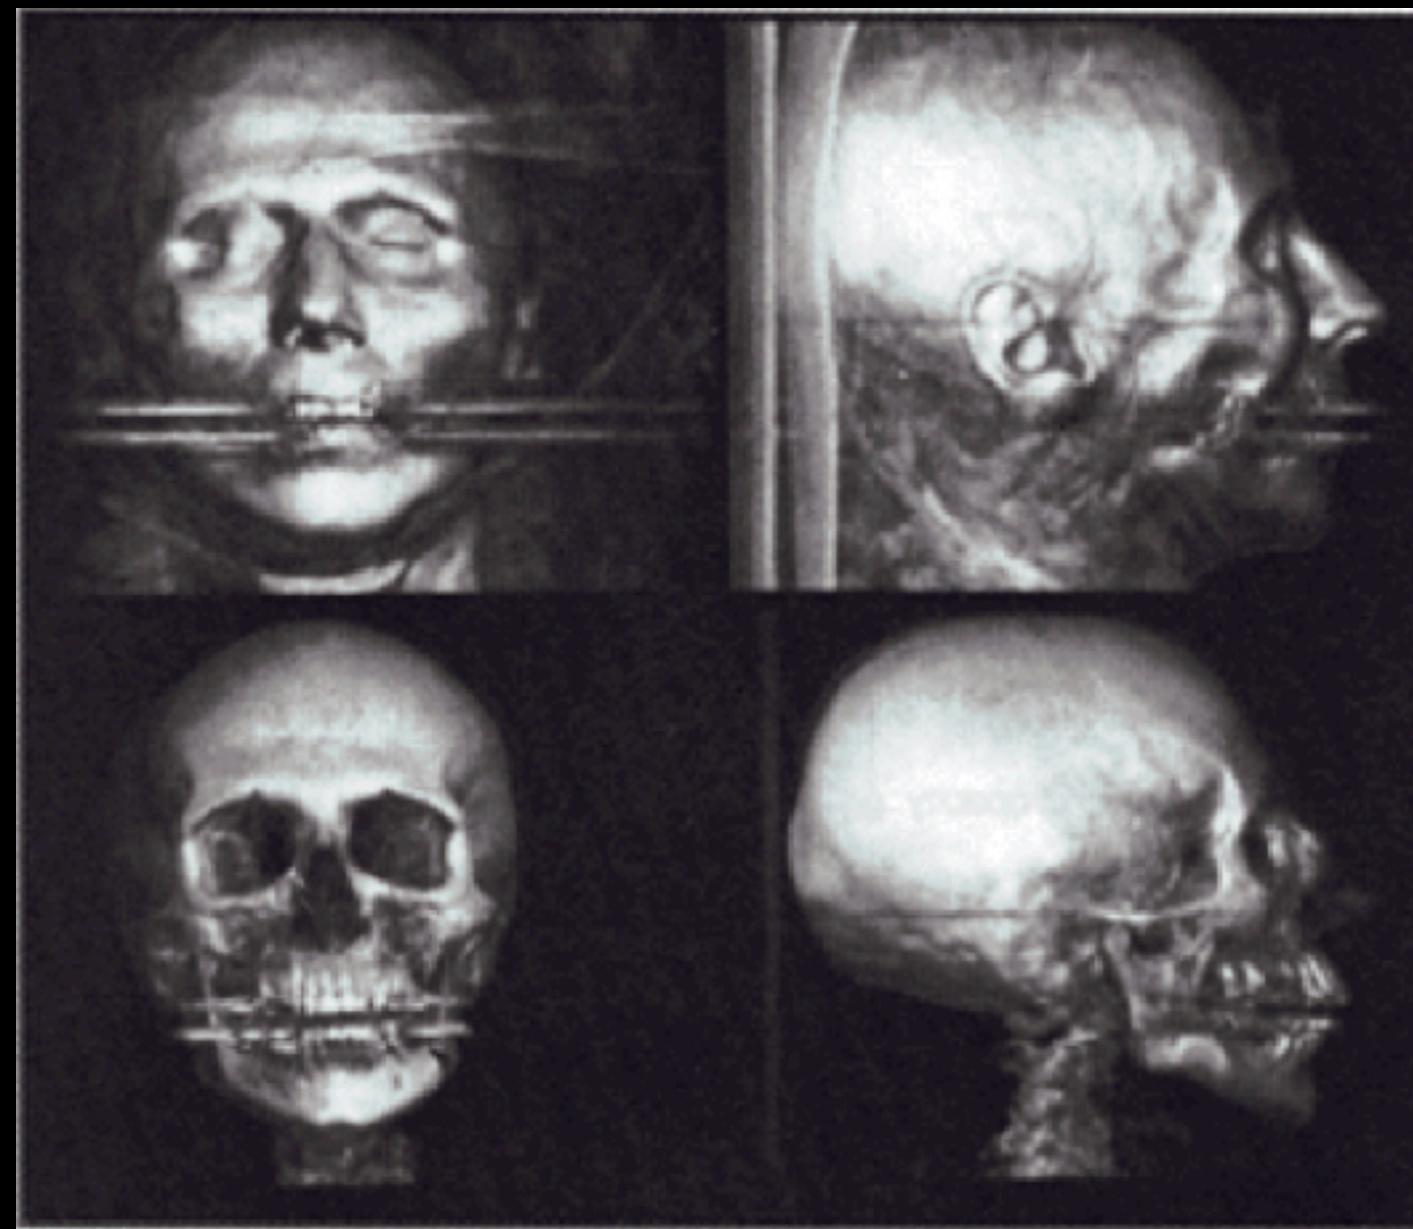


ALEXANDER BOCK
MOORE-SLOAN POSTDOCTORAL FELLOW
CENTER FOR DATA SCIENCE
NEW YORK UNIVERSITY

- Slides:
 - https://github.com/alexanderbock/alexanderbock.github.io/blob/master/lectures/2018/ds_ga_3001_017_volumerendering.pdf
 - alexanderbock.eu/lectures/2018/ds_ga_3001_017_volumerendering/index.html

DIRECT VOLUME RENDERING

DIRECT VOLUME RENDERING



Levoy, 1988

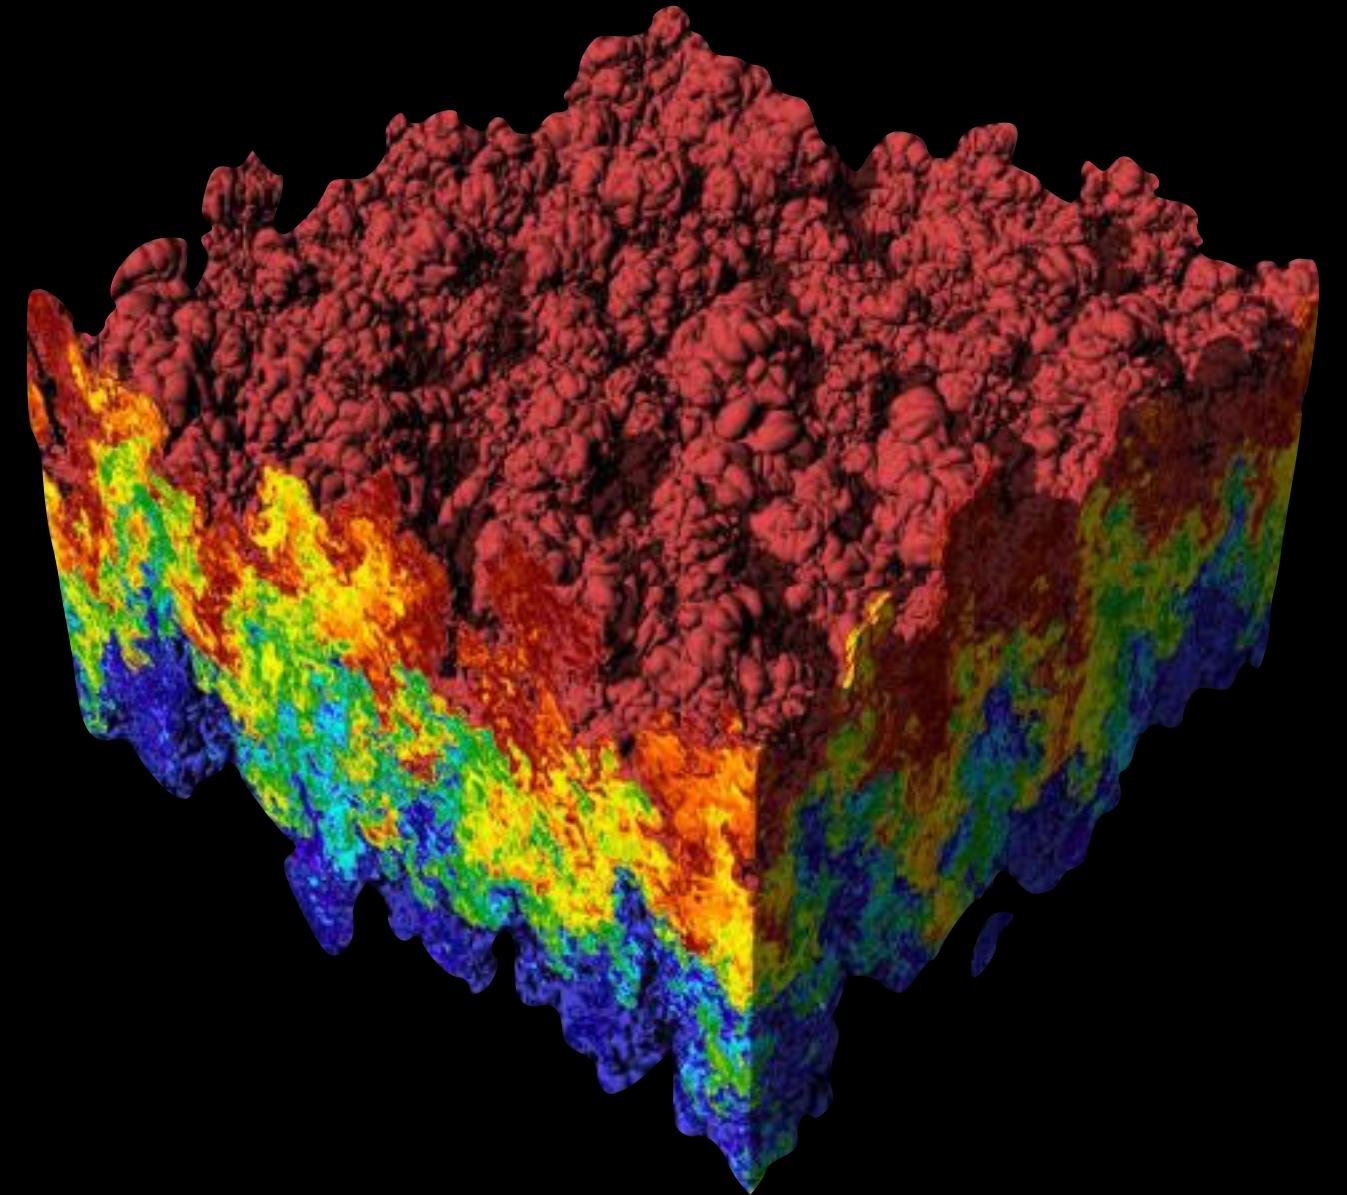
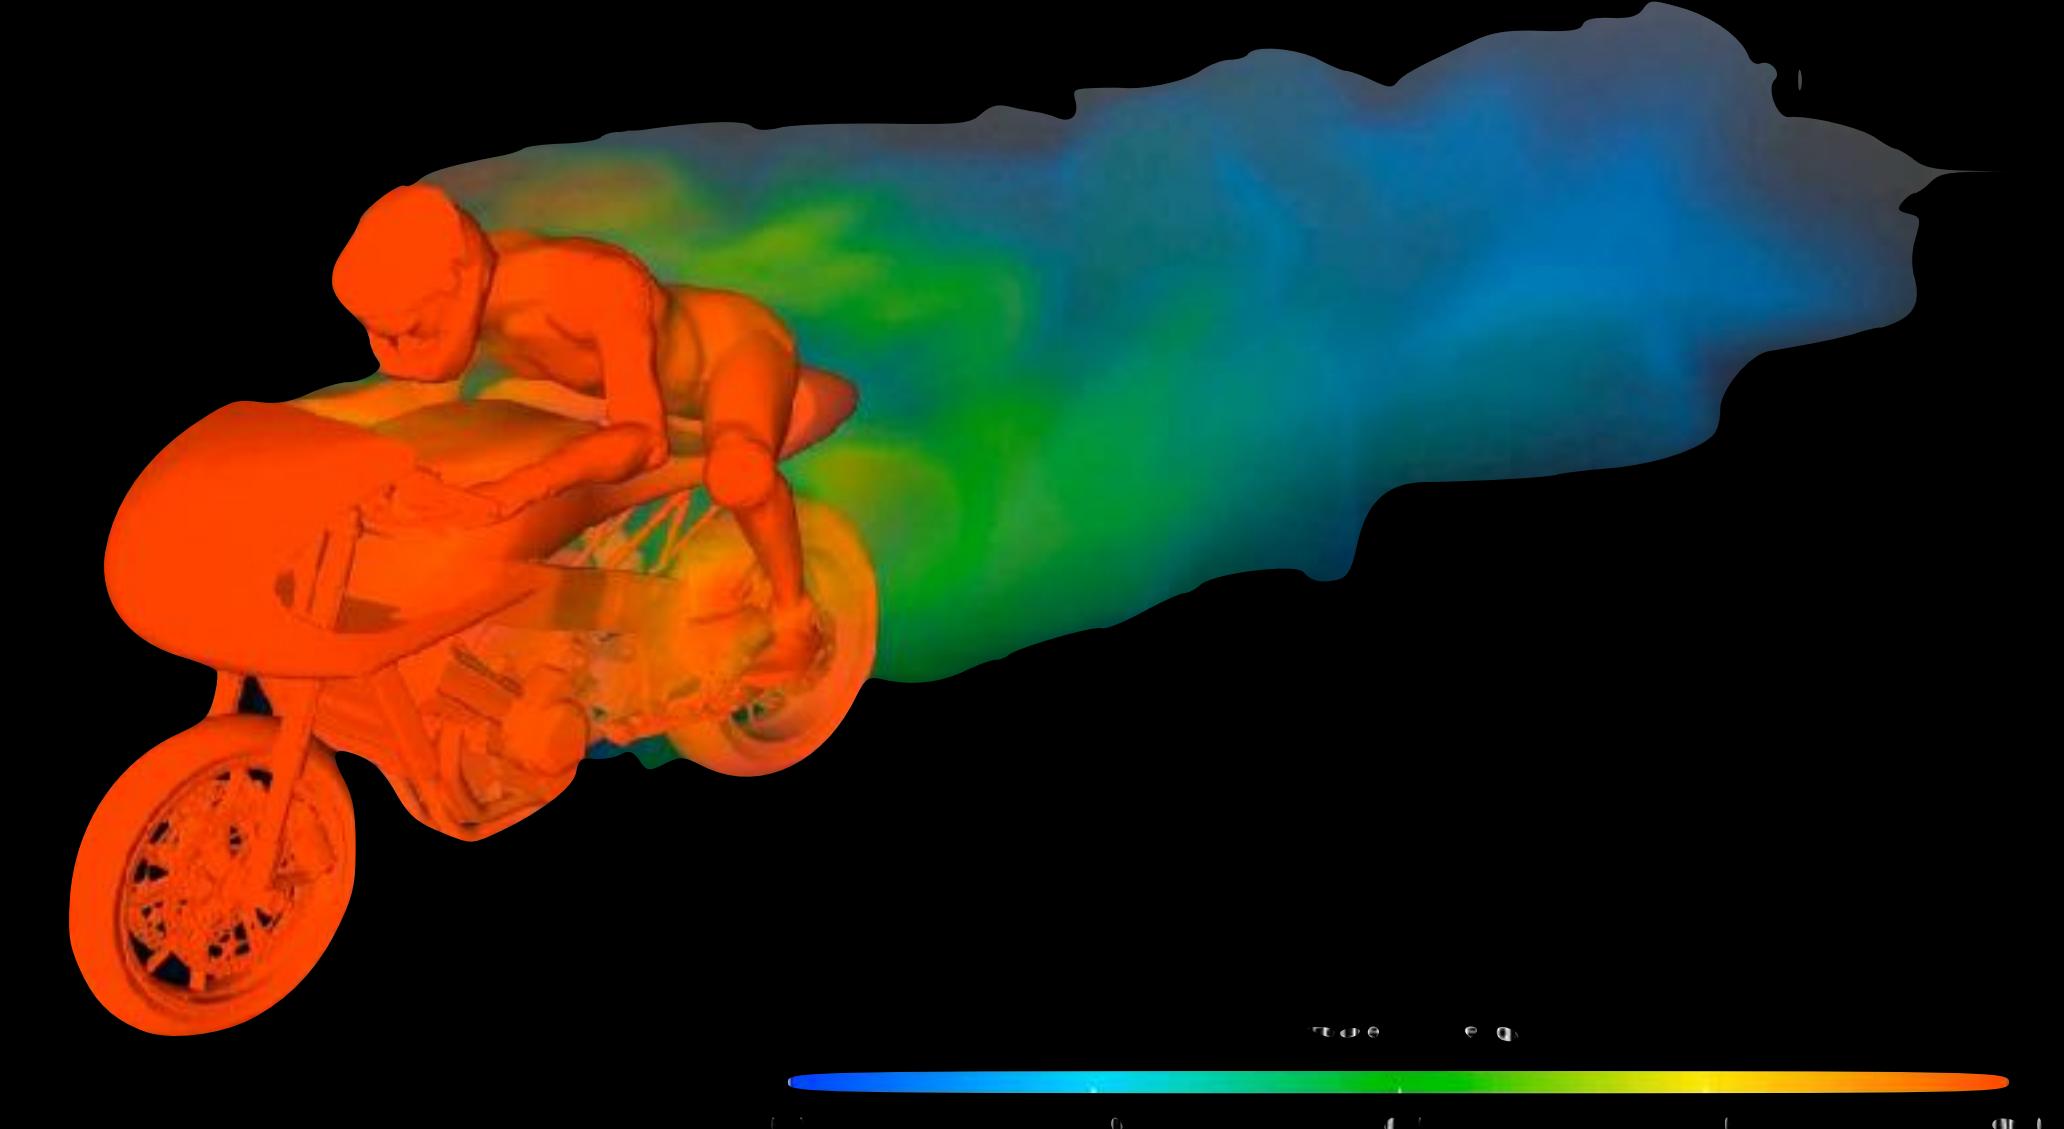
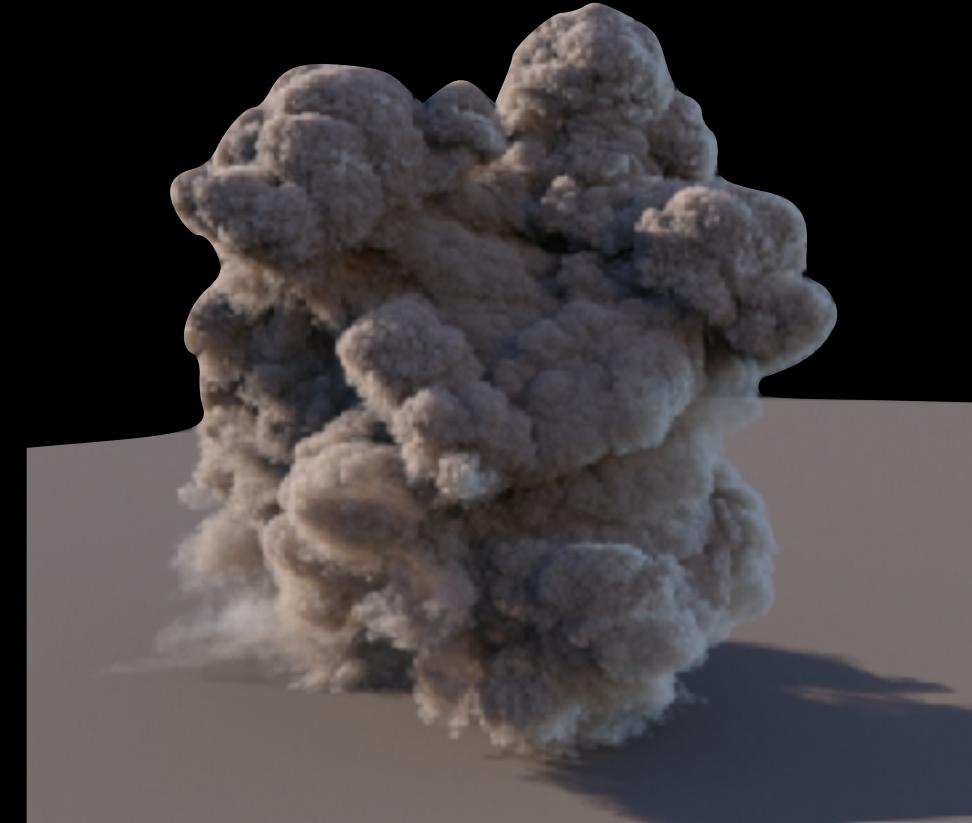


Kronander, 2012



Jönsson, 2012

DIRECT VOLUME RENDERING



DIRECT VOLUME RENDERING

- **Direct**
- No explicit geometric representation is generated
- **Volume**
- Extension of 2D image into 3D
- For today: Regular grid structure, but other methods exist (AMR, spherical, ...)
- **Rendering**
- ...

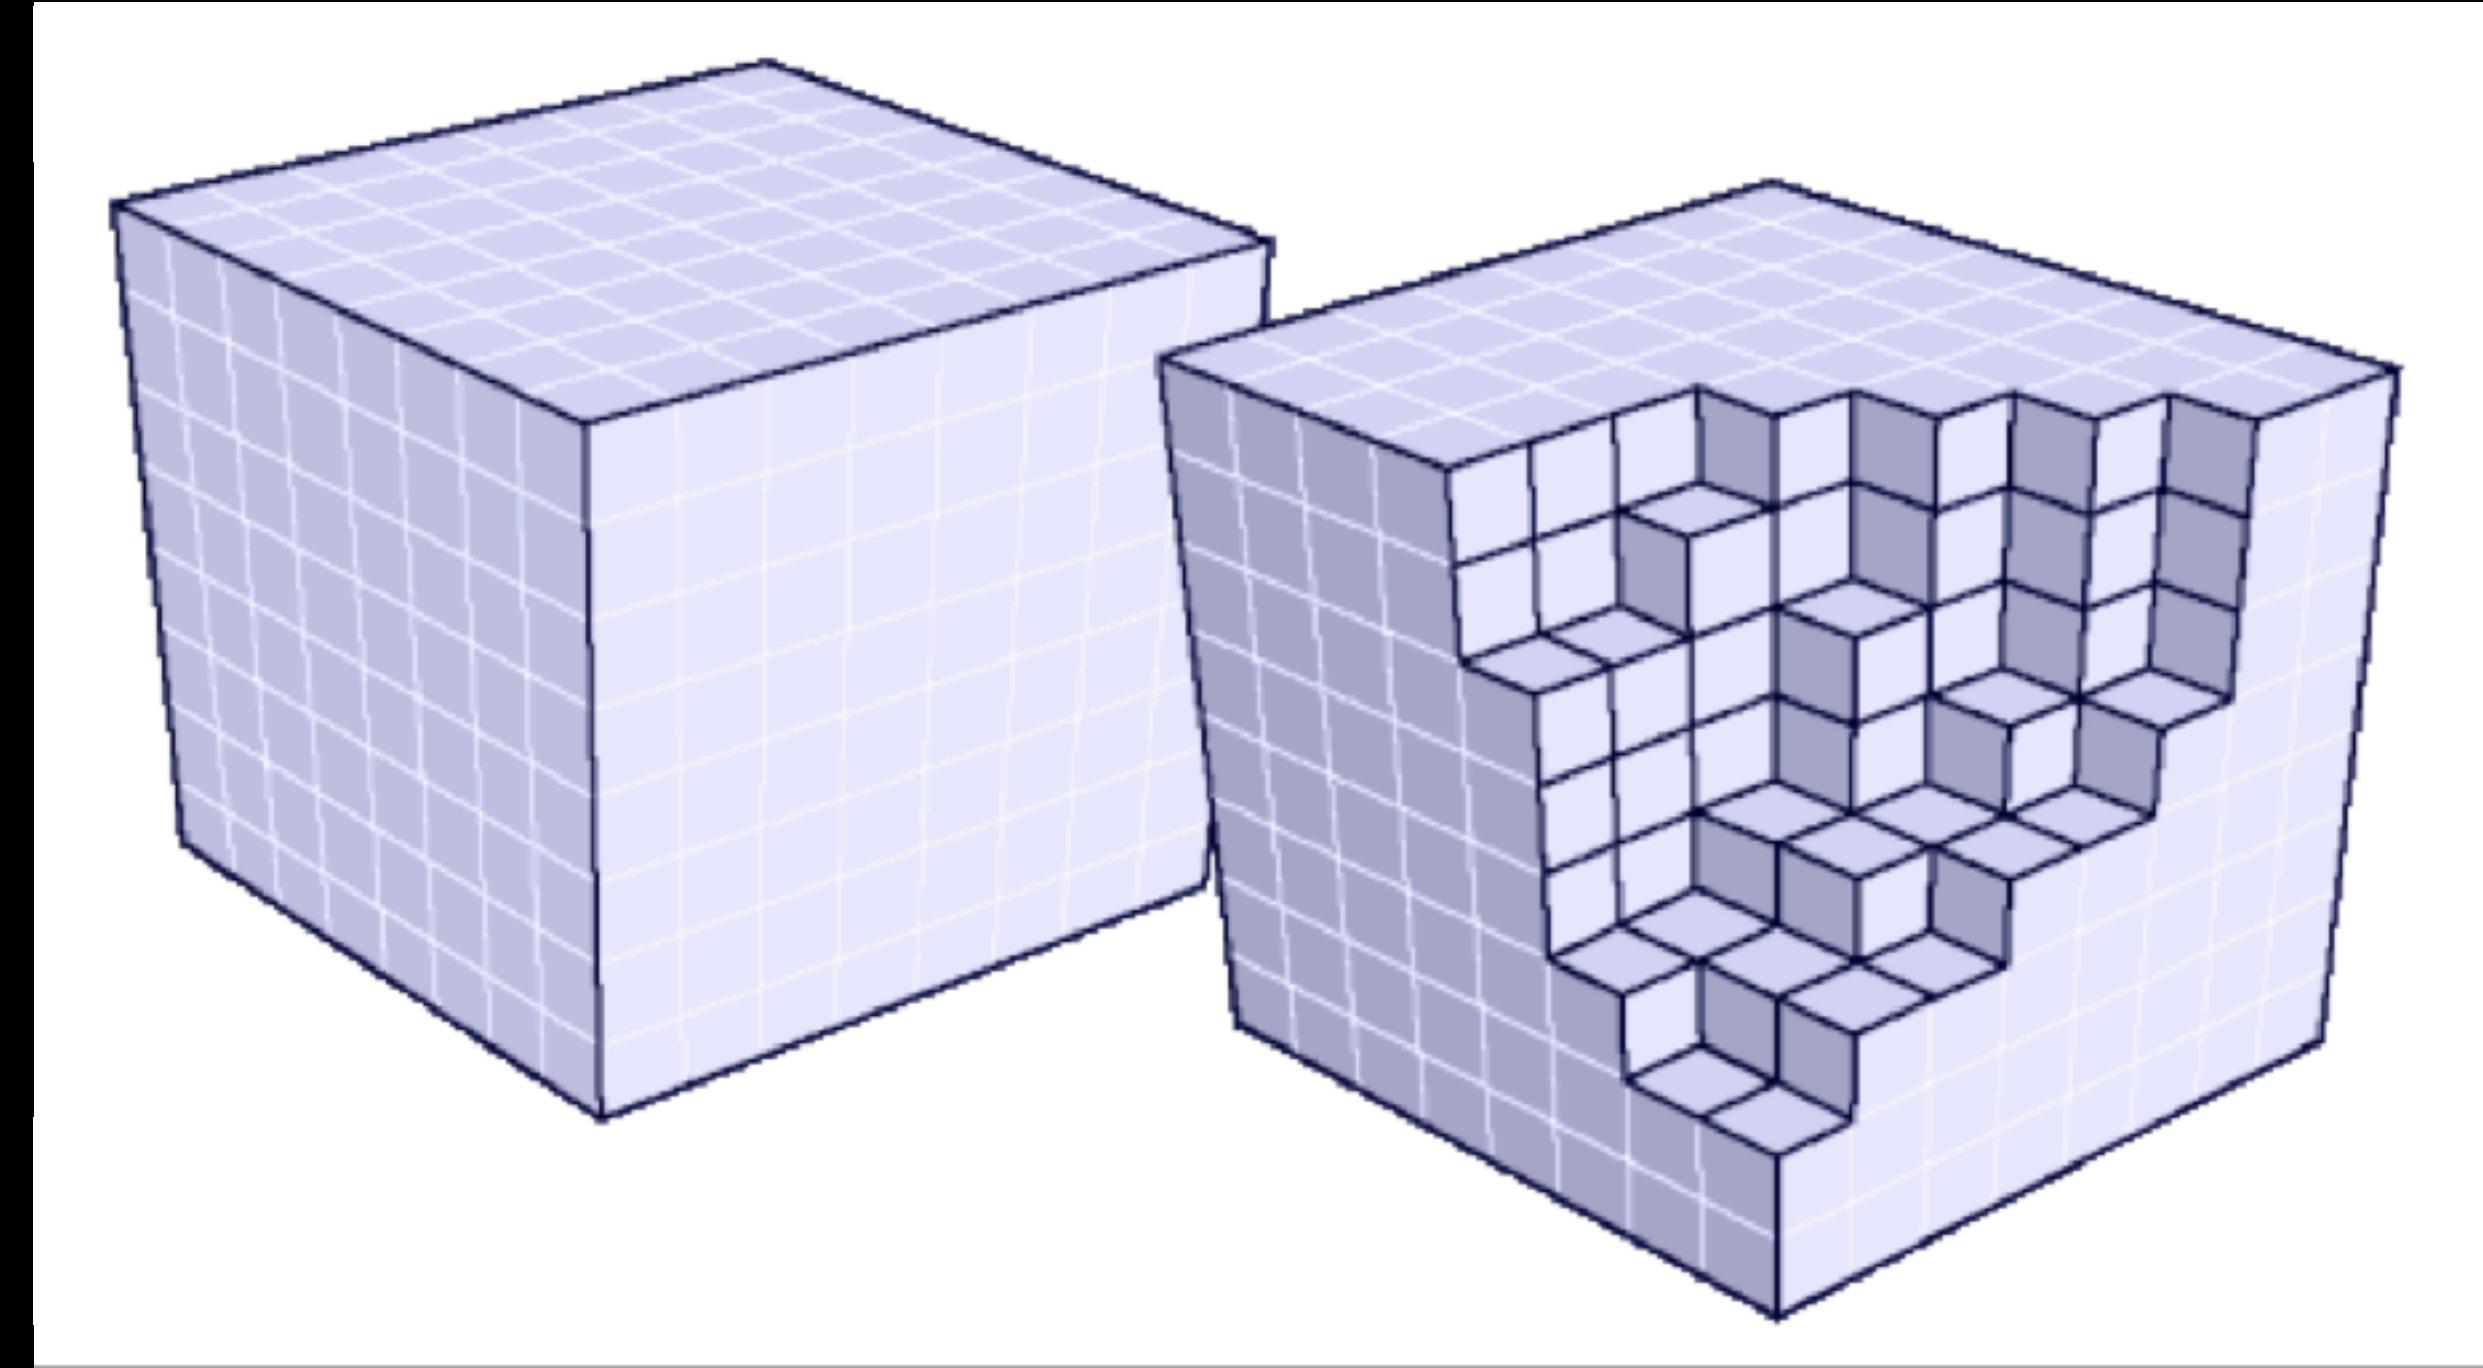


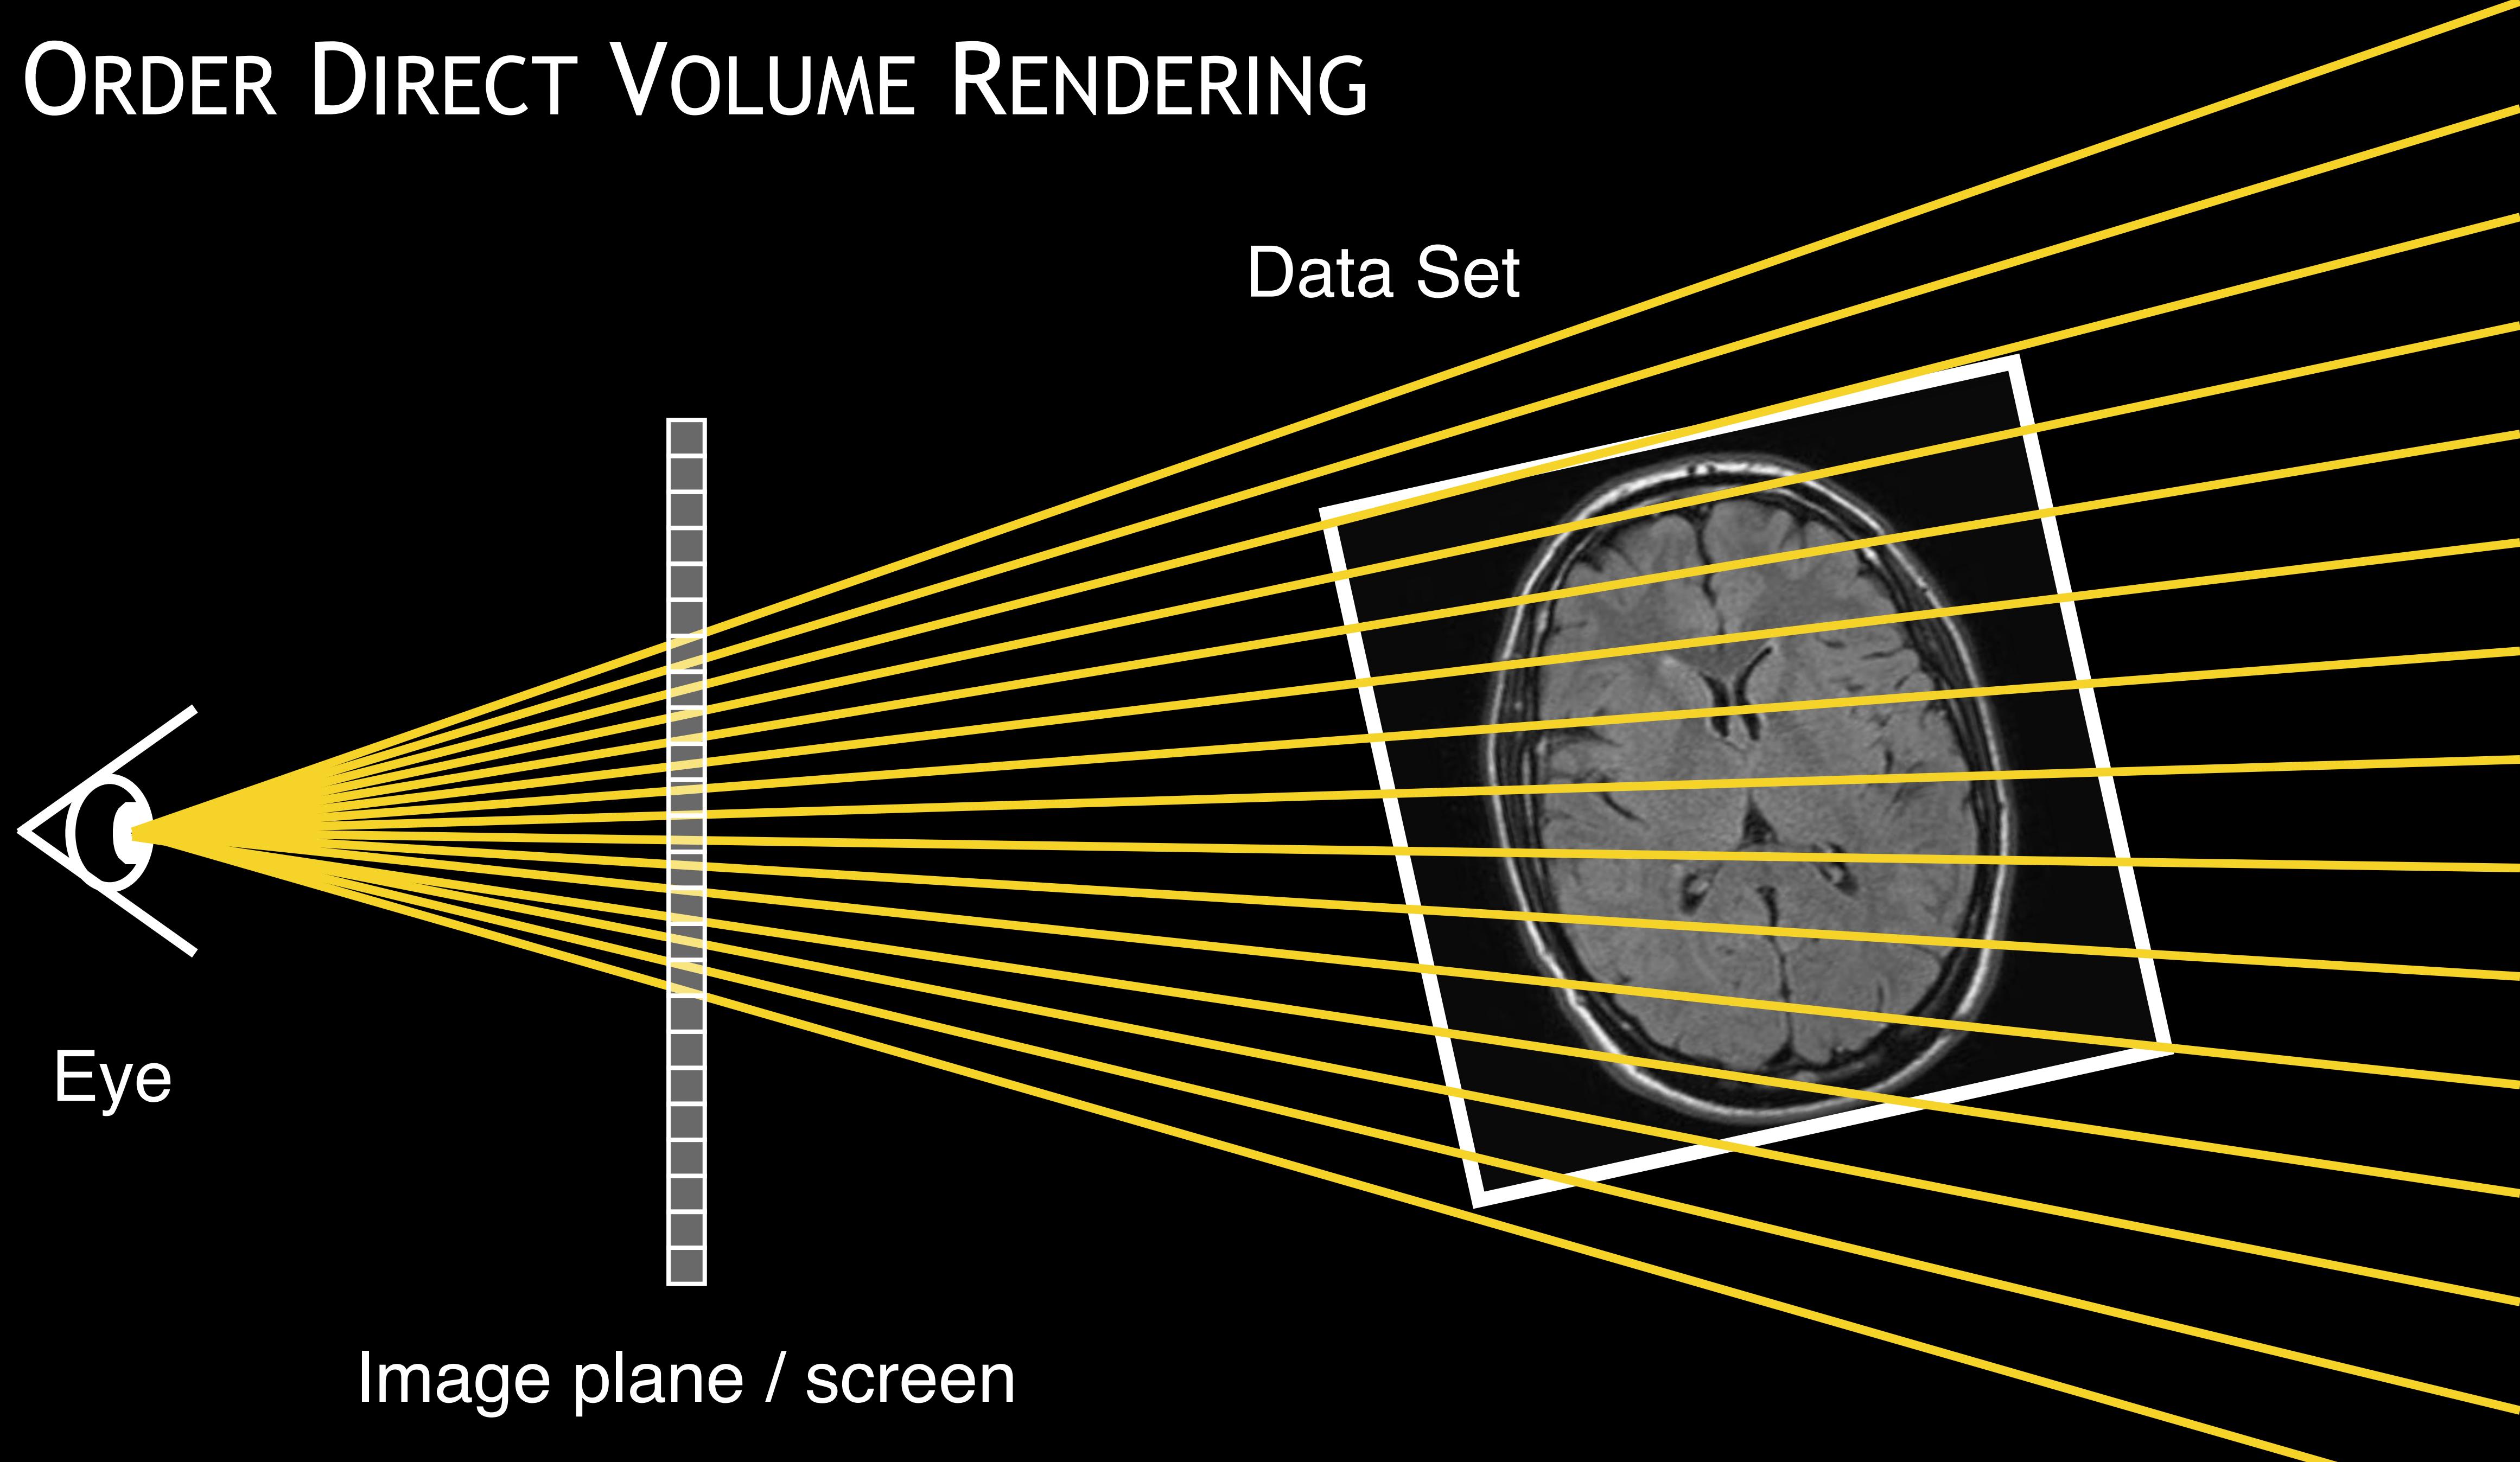
Image © Weiskopf, Machiraju, Möller

DIRECT VOLUME RENDERING

- Image-order techniques
 - Perform operations on pixels to produce the image
 - -> Majority of today's implementations
- Object-order techniques
 - Perform operations on the elements of the scene
 - -> What we have been doing so far

Image © Weiskopf, Machiraju, Möller

IMAGE ORDER DIRECT VOLUME RENDERING



EMISSION ABSORPTION MODEL

- Assumption

Volume consists of small particles which are

- opaque
- non-reflecting
- light emitting
- light absorbing
- the only light sources in the scene

RENDERING INTEGRAL

$$I(\vec{x}_c) = I_0(\vec{x}_0) T(\vec{x}_0, \vec{x}_c) + \int_{\vec{x}_c}^{\vec{x}_0} T(\vec{x}_c, \vec{x}) \sigma_\alpha(\vec{x}) I_c(\vec{x}) d\vec{x}$$

- No closed form solution of the integral in general
- Numerical approach to compute volume rendering integral:

Riemann Sum

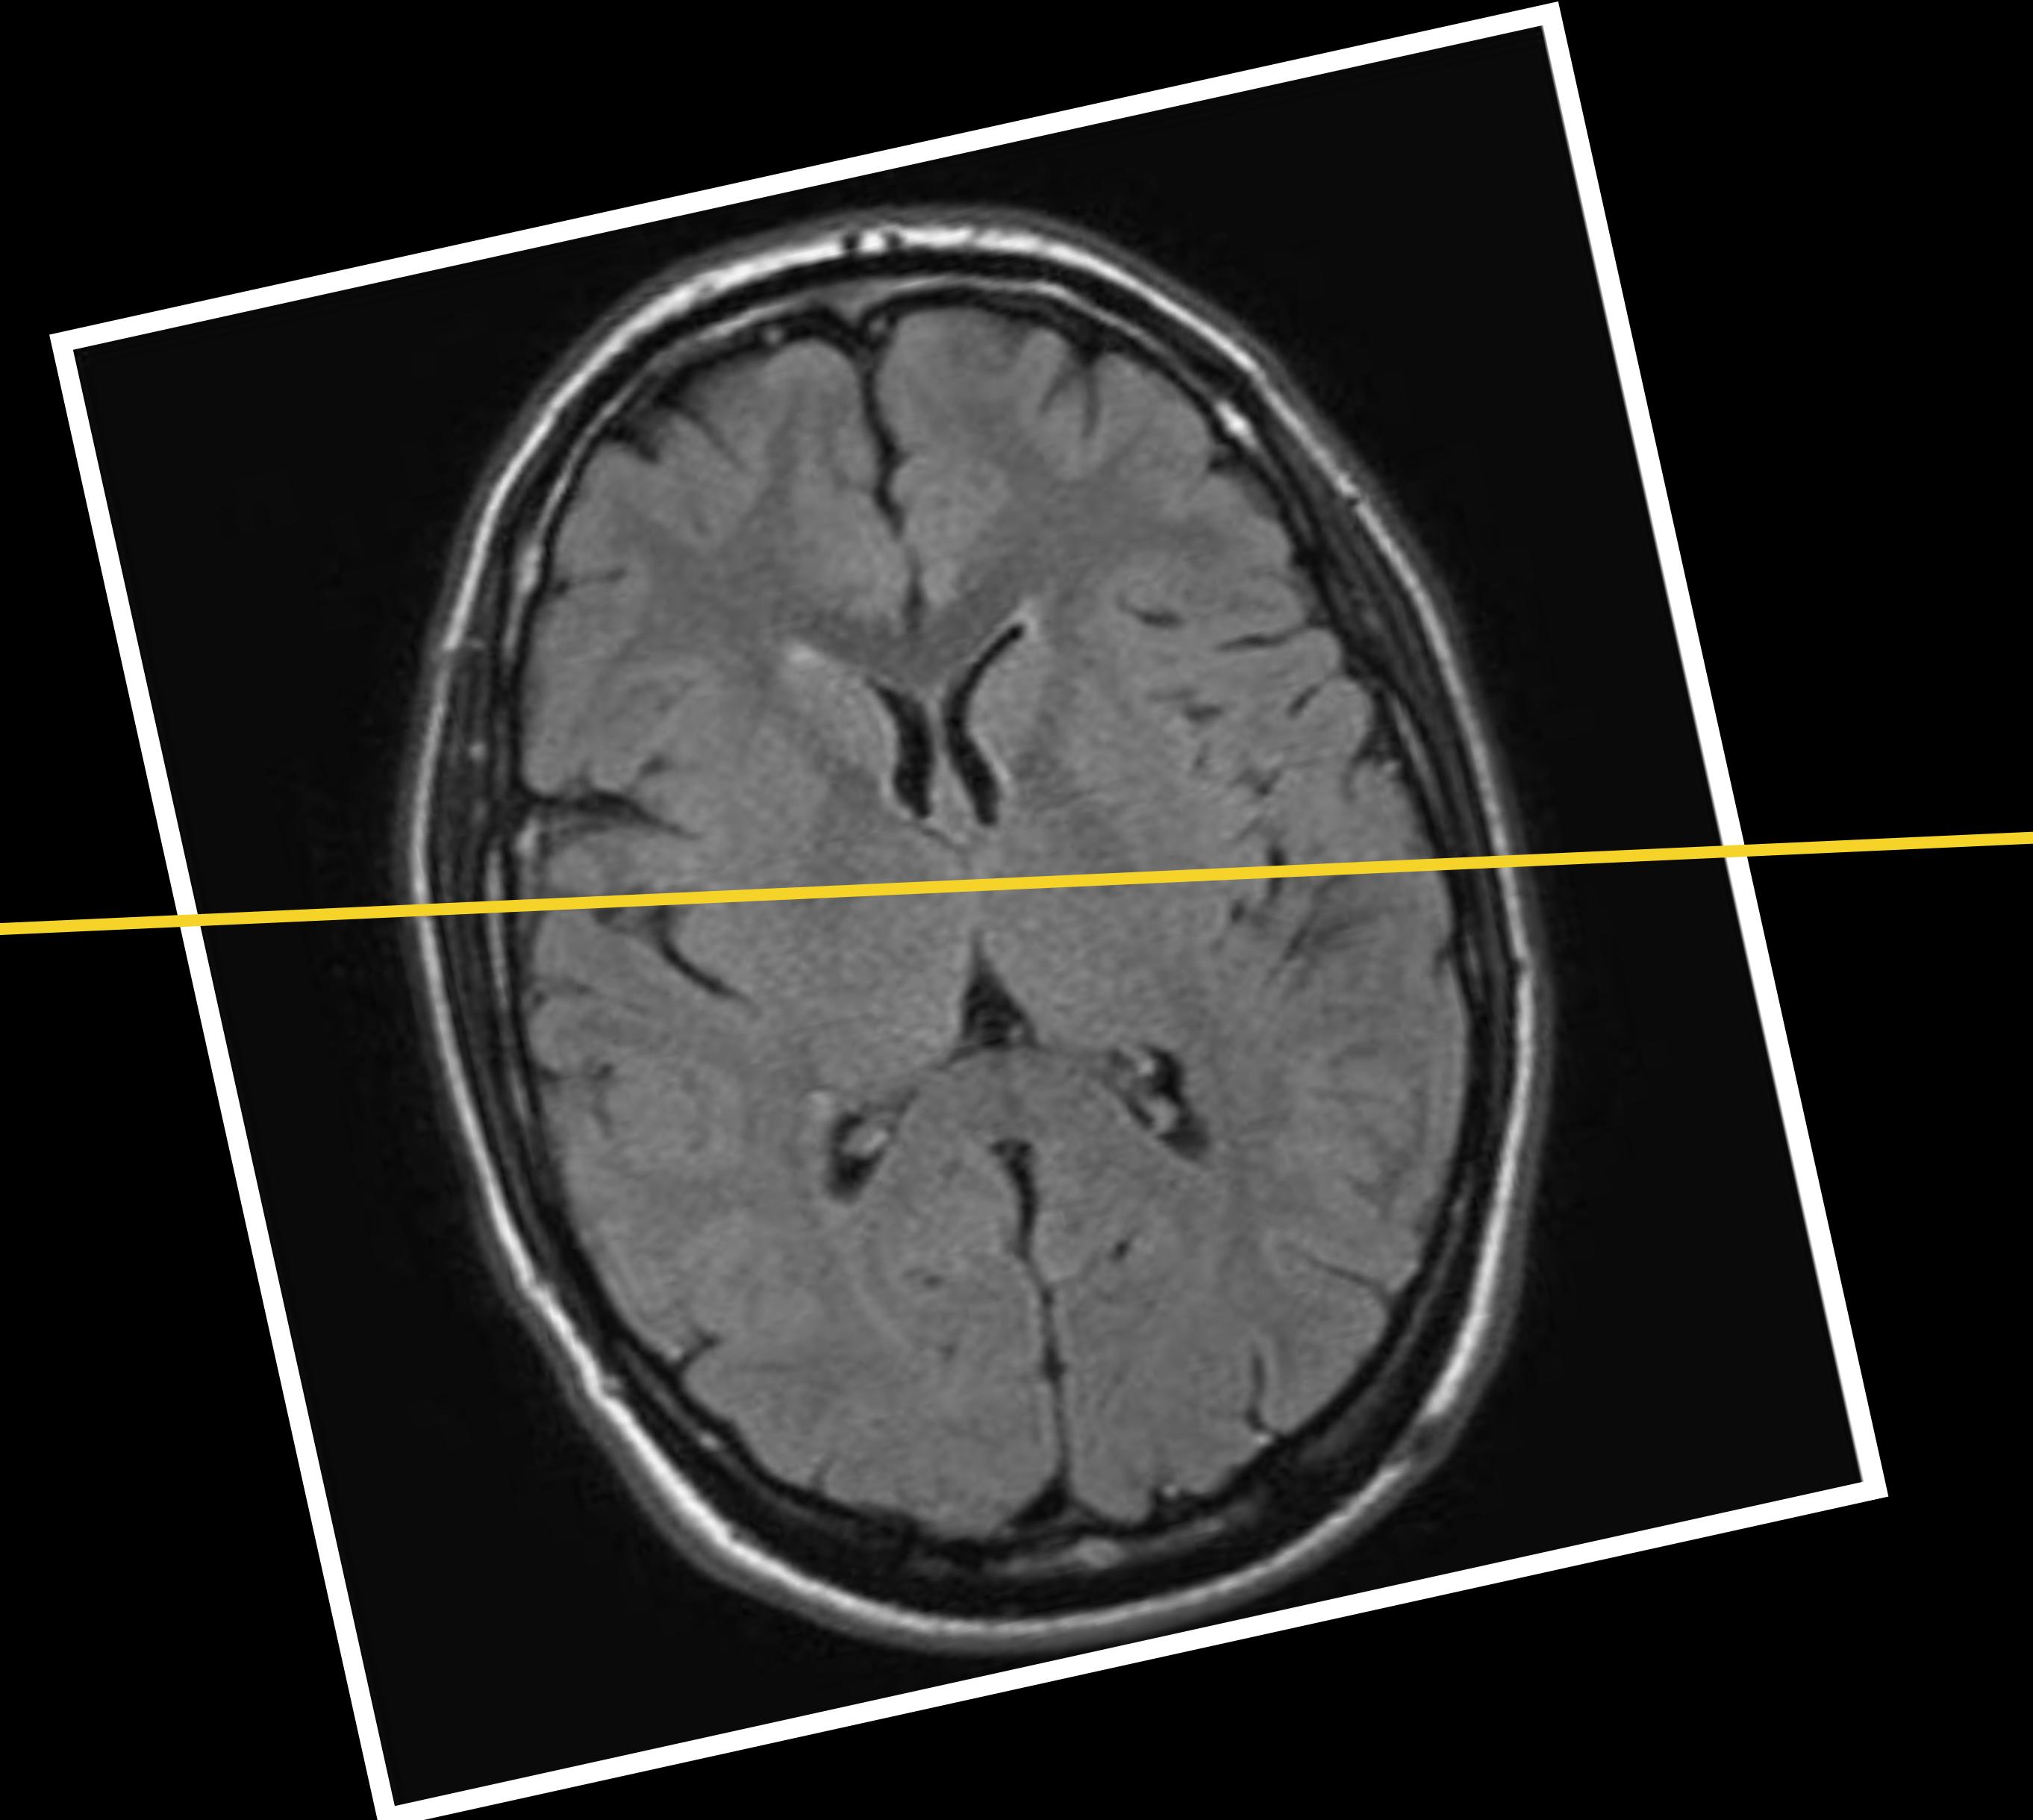
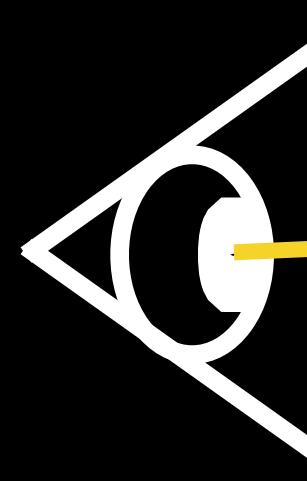
$$T(\vec{a}, \vec{b}) = \exp \left(- \int_{\vec{a}}^{\vec{b}} \tau(\vec{x}) d\vec{x} \right)$$

I : Intensity
 τ : Absorption factor
 σ : Emission / Absorption

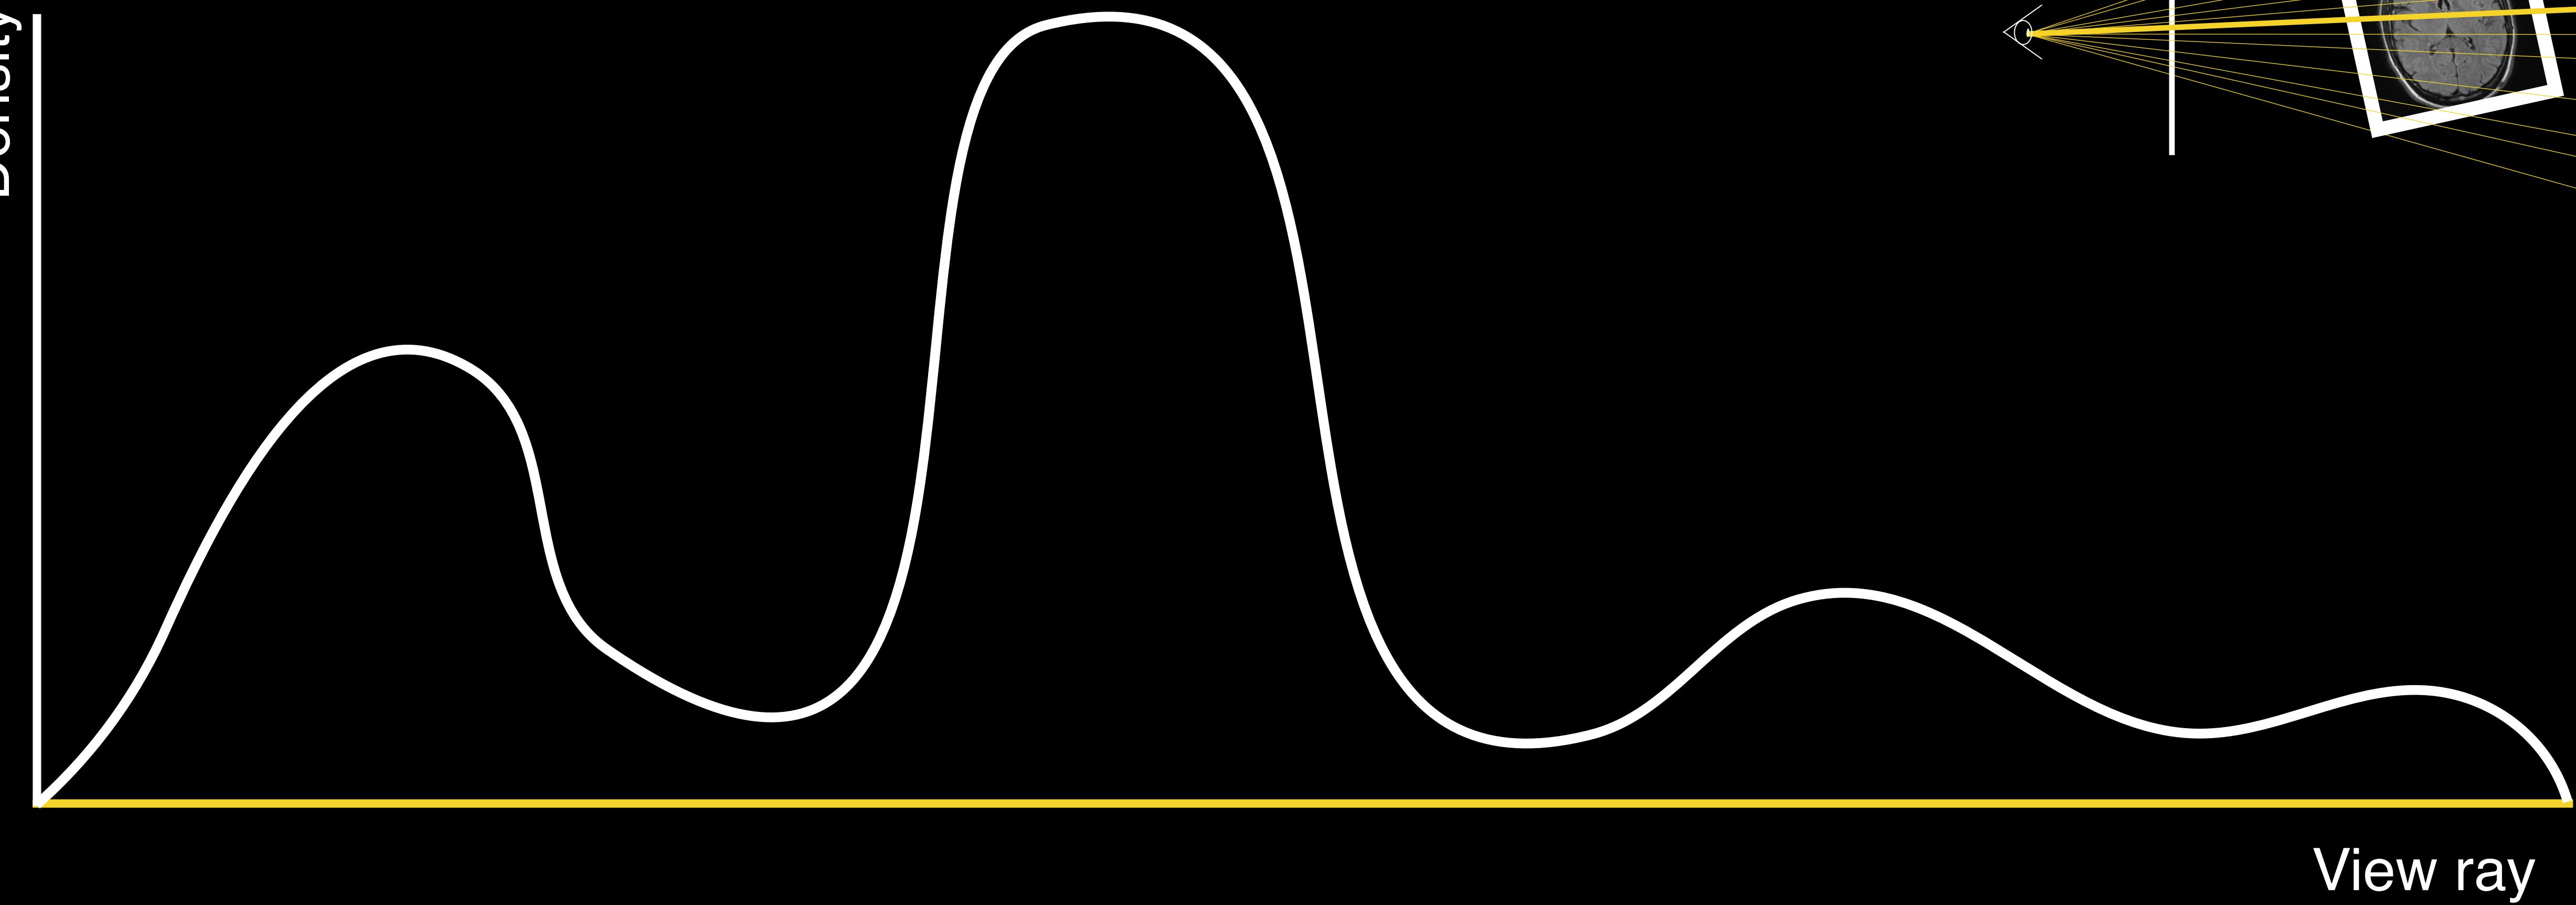
Max, Optical Models for Direct Volume Rendering, 1995

RAYCASTING

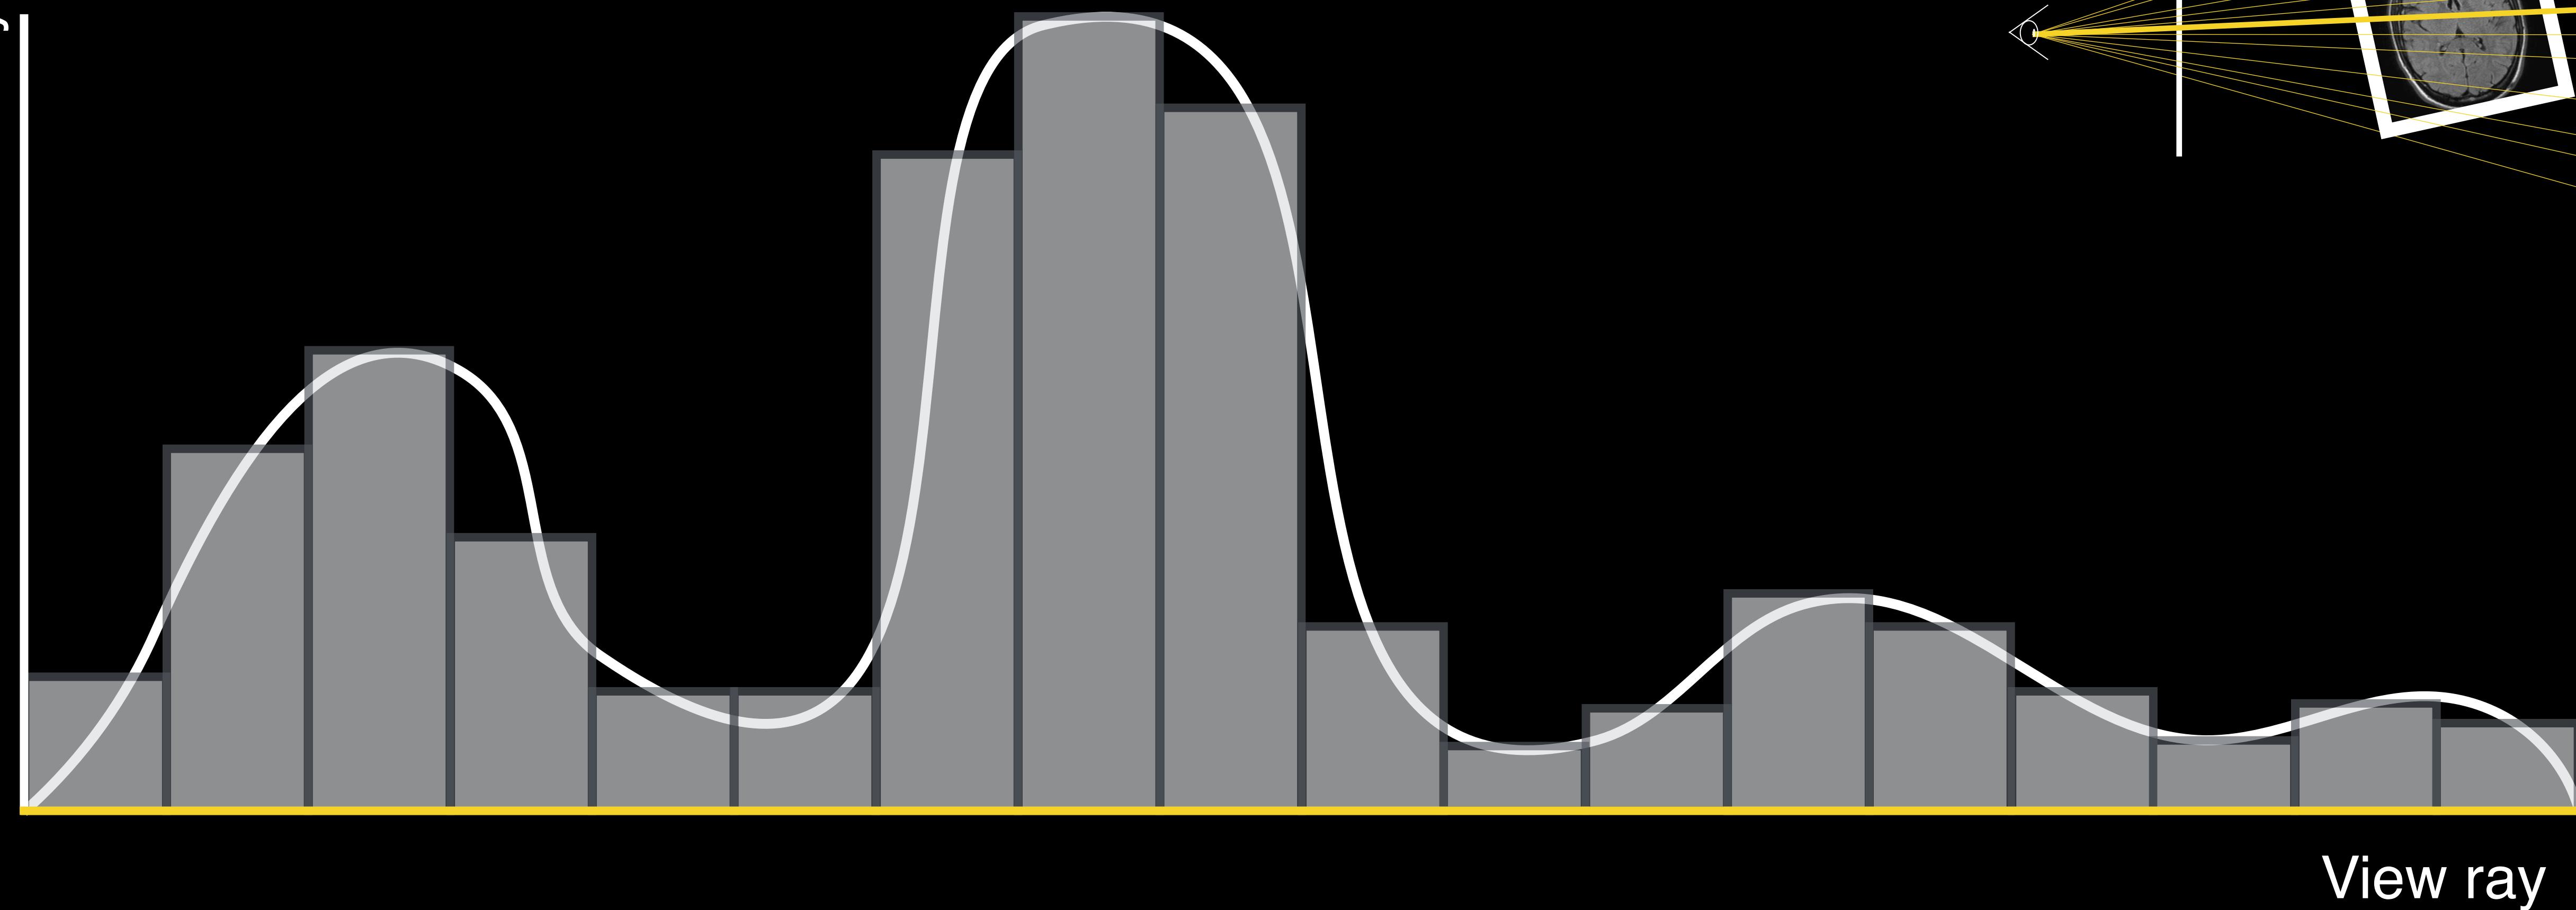
- Focus on a single viewing ray first



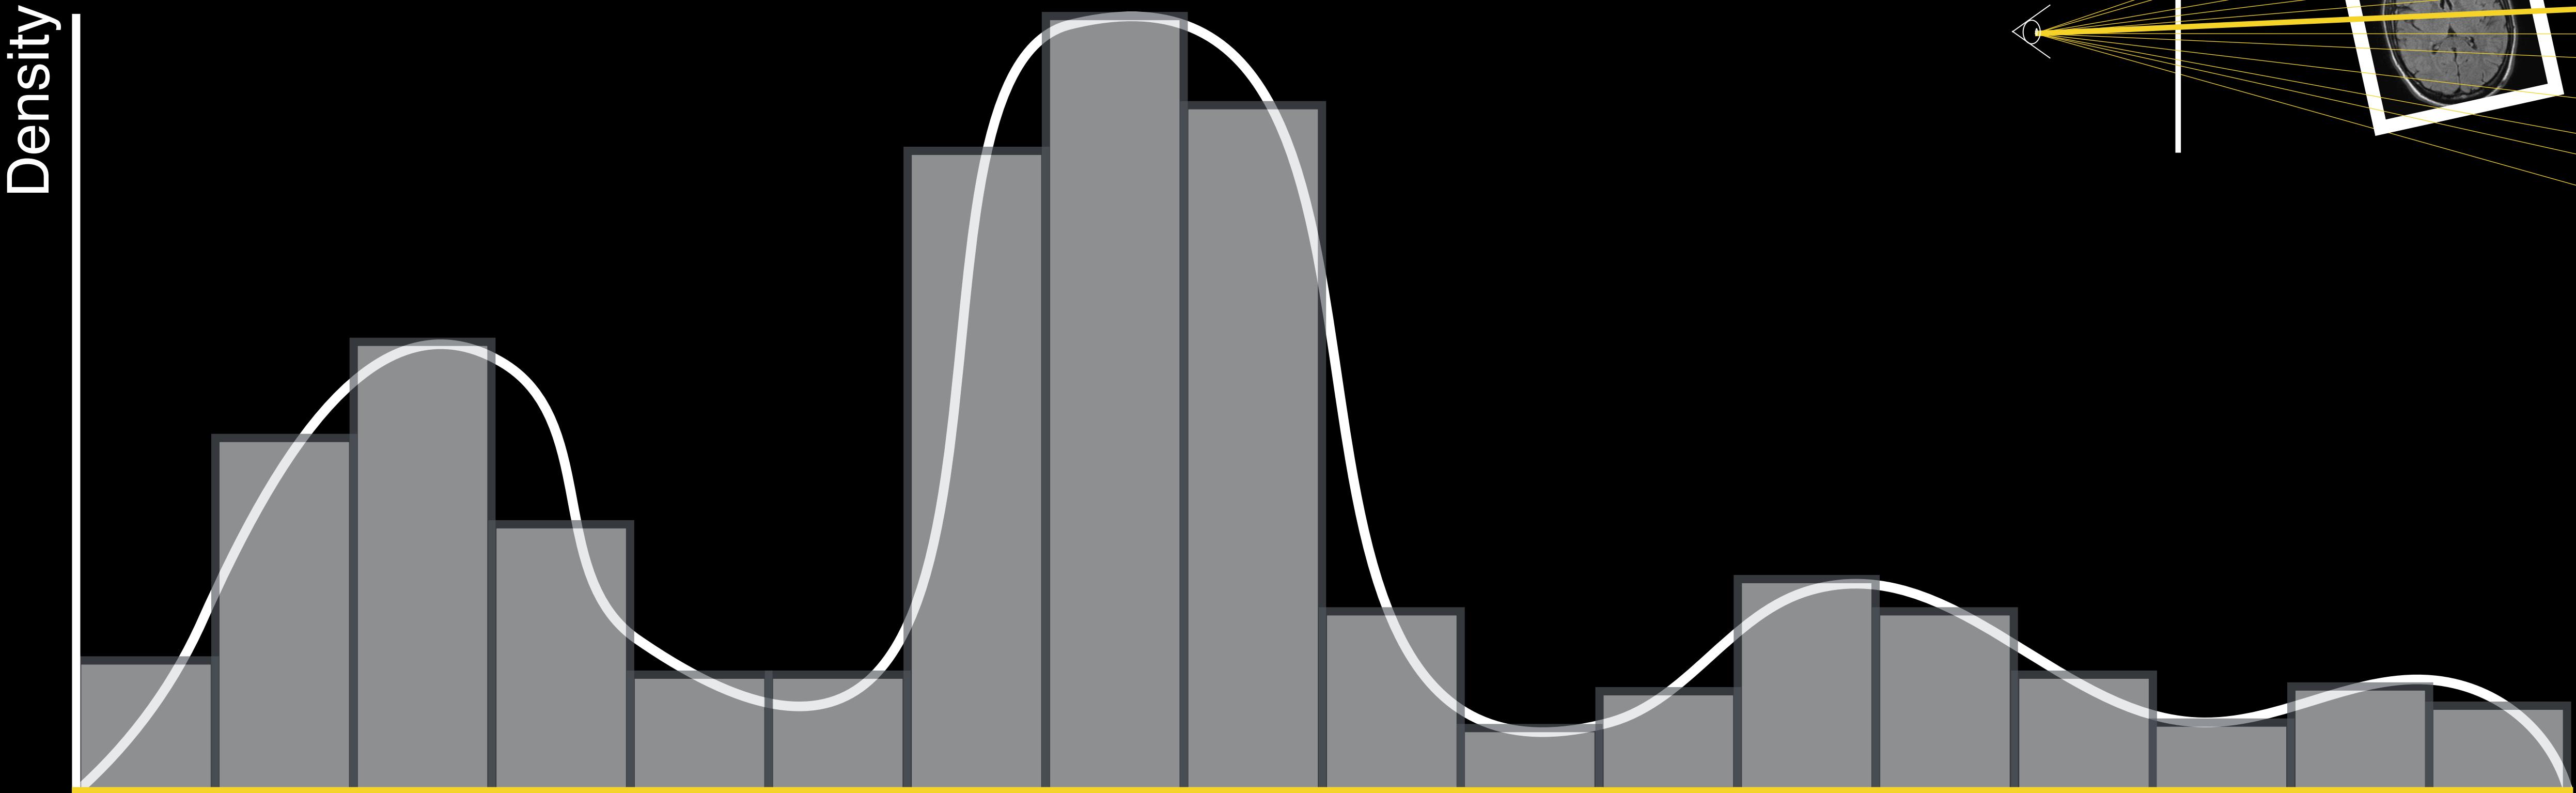
RAYCASTING



RAYCASTING



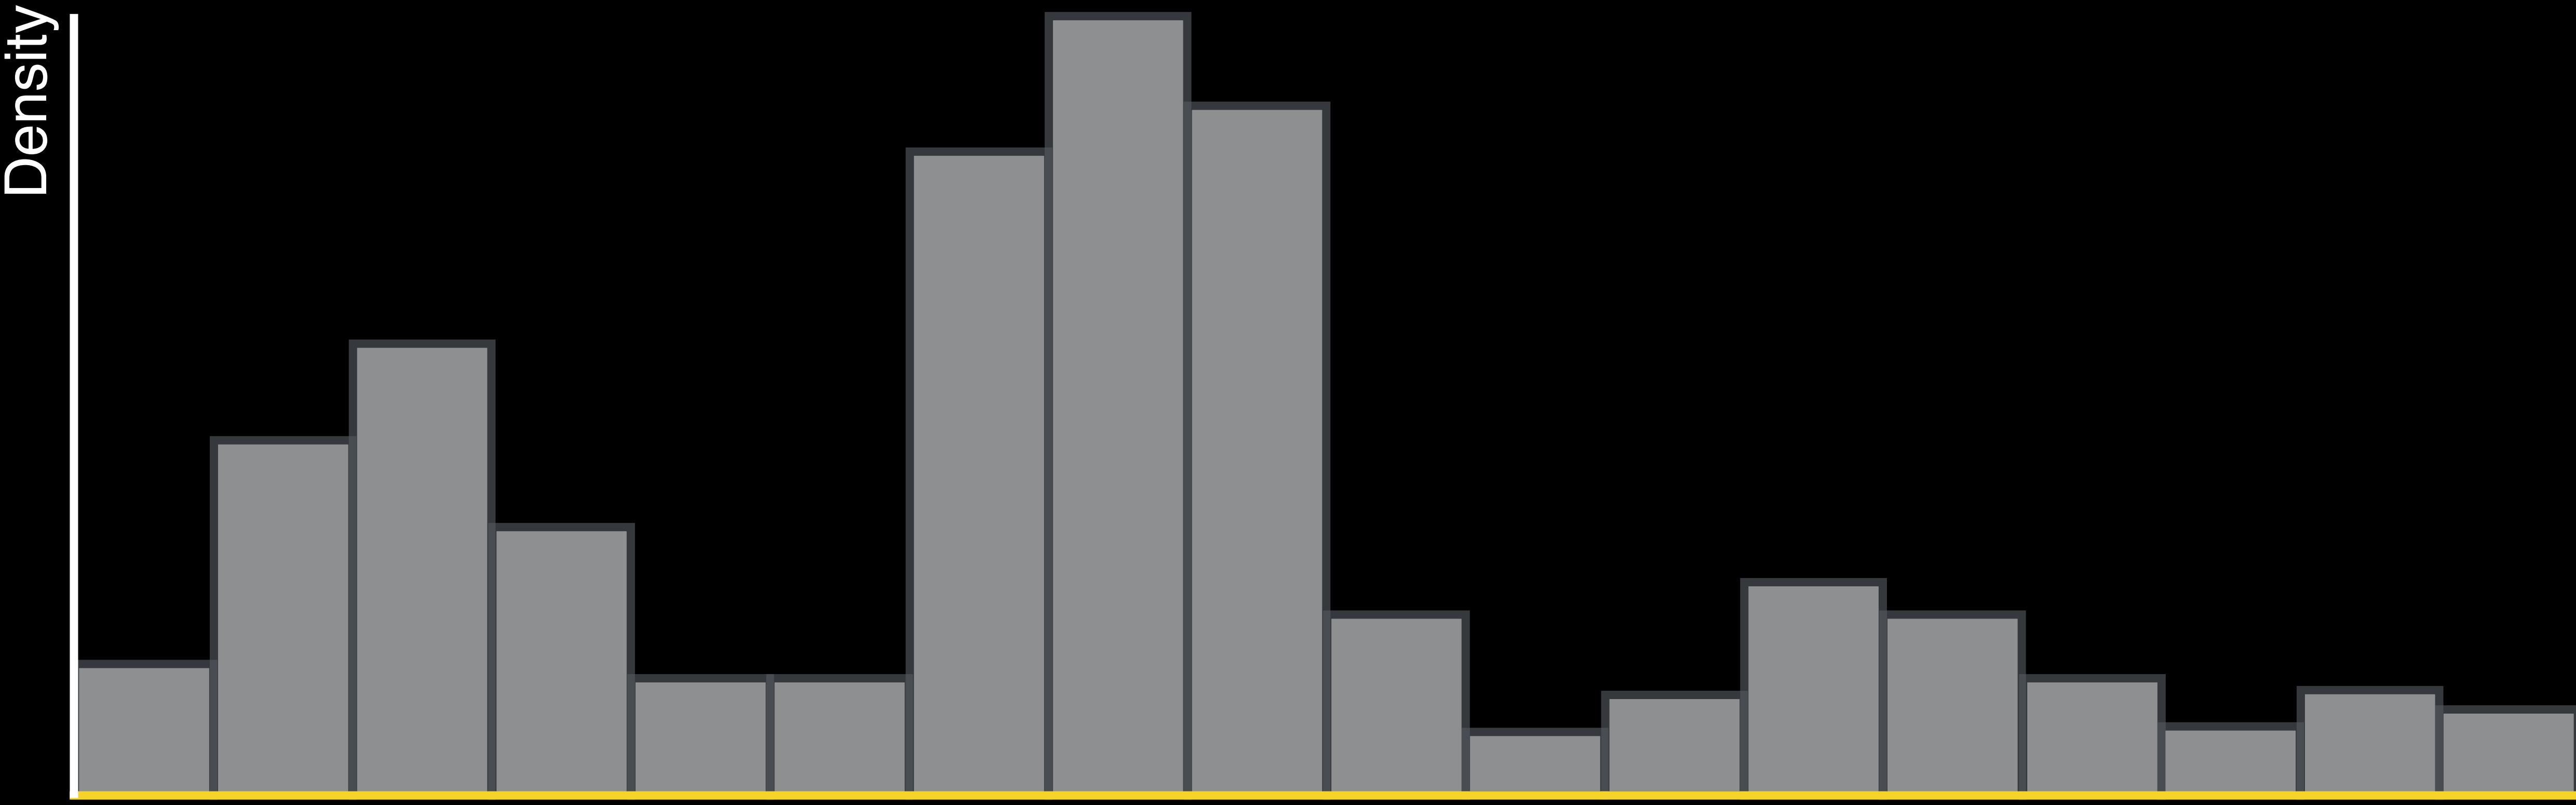
RAYCASTING



- Nyquist-Shannon sampling theorem
 - Required sampling frequency: $2 * \text{highest frequency}$ to completely reconstruct a signal
 - Problem: Frequency of the dataset is, in general, unknown -> user parameter

COMPOSITING

COMPOSITING



- Front-to-back
- Maximum intensity projection
- First Hit

View ray

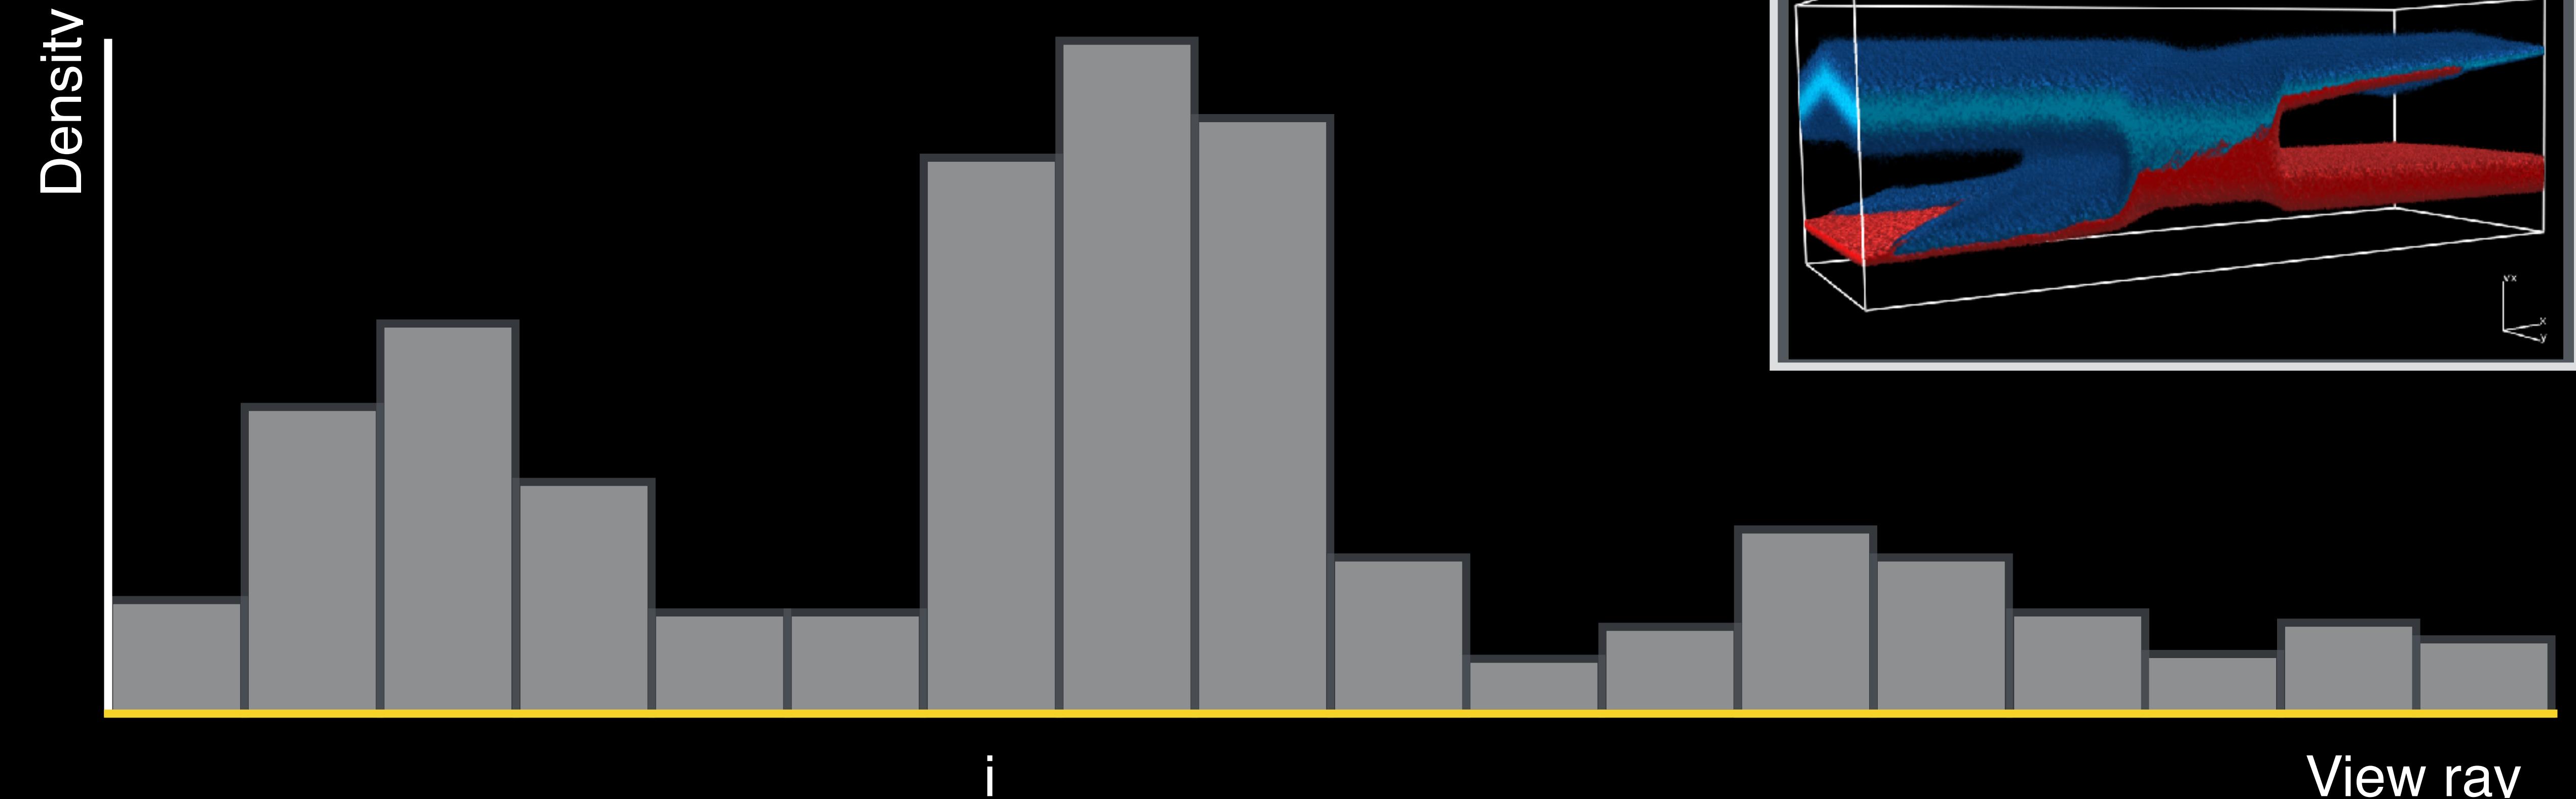
C: Color

A: Transparency/Alpha

Prime: Accumulated

Non-prime: Current value

FRONT-TO-BACK COMPOSITING



- See OpenGL `glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA)`
- $C'_i = (1 - A'_{i-1})C_i + C'_{i-1}$
- $A'_i = (1 - A'_{i-1}) * A_i + A'_{i-1}$
- Stop when full opacity has been reached (or *Early Ray Termination*)

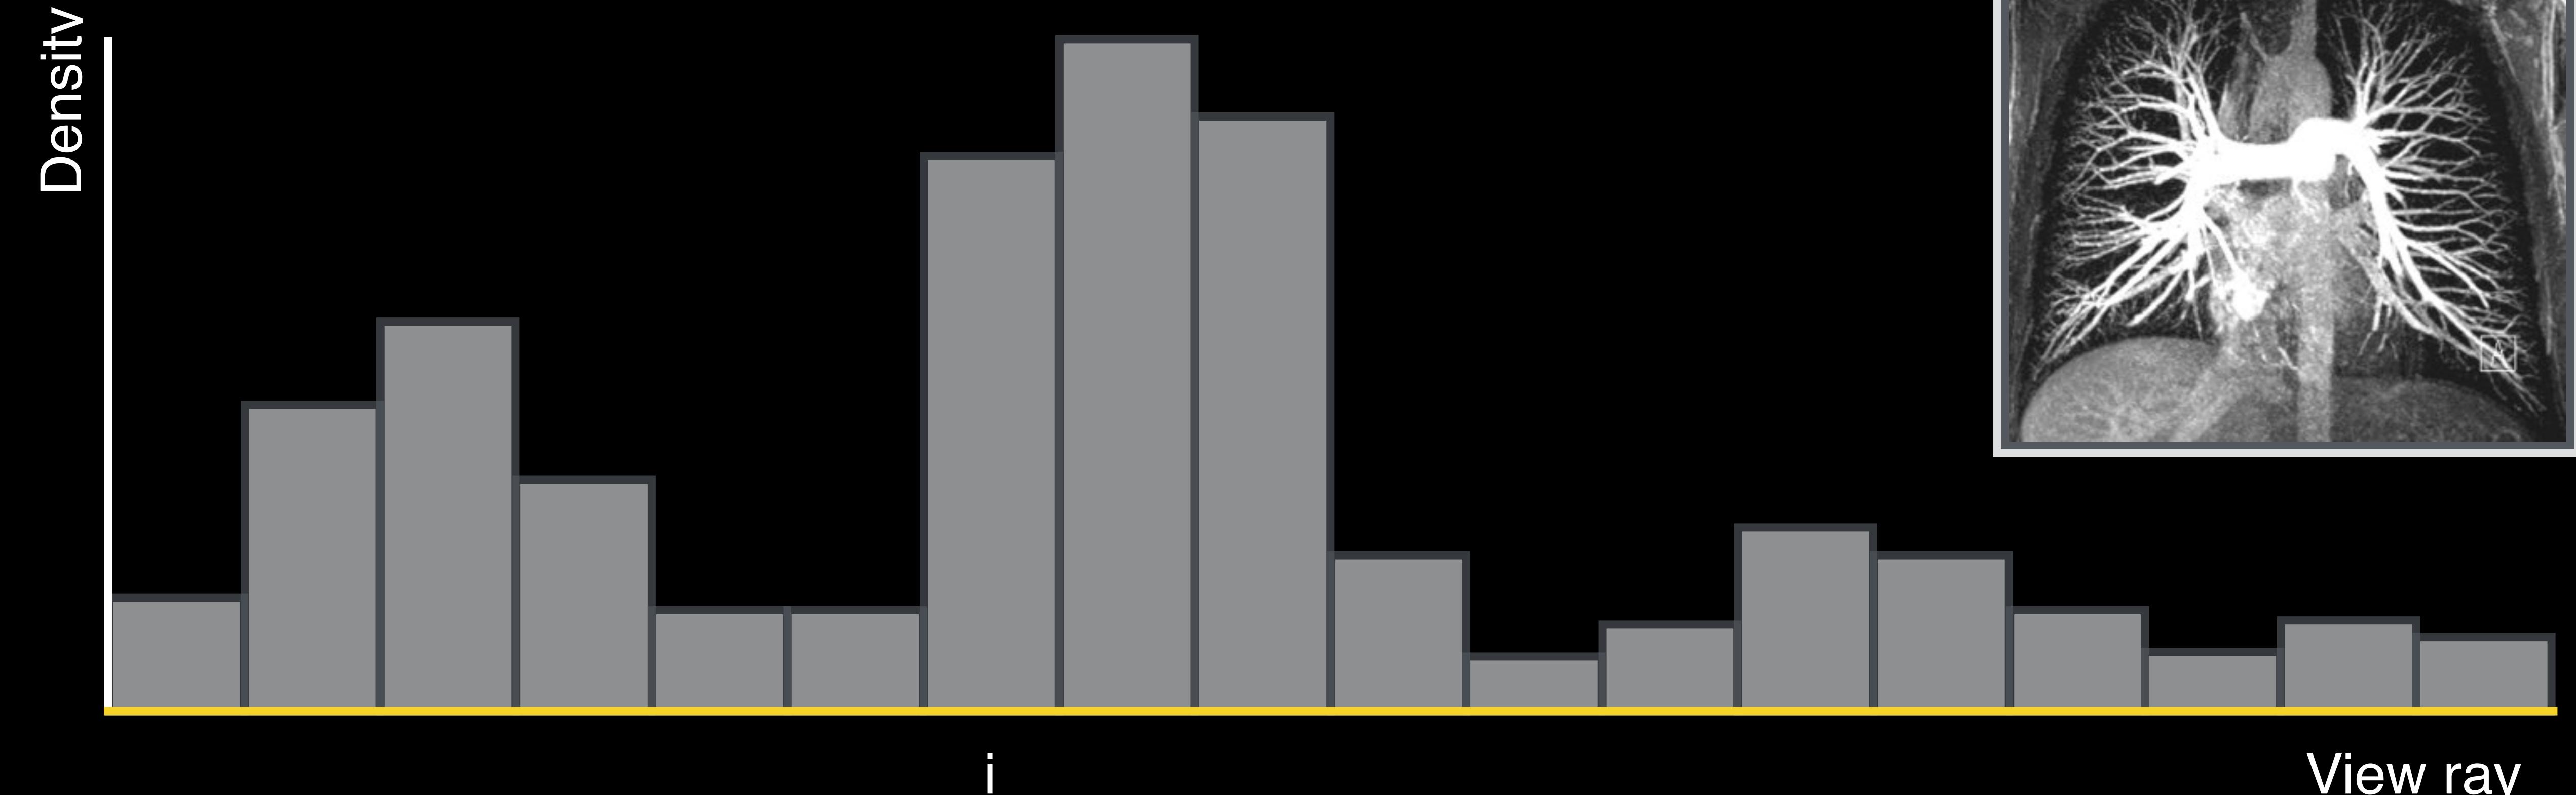
C: Color

A: Transparency/Alpha

Prime: Accumulated

Non-prime: Current value

MAXIMUM INTENSITY COMPOSITING



- Use the maximum intensity of all samples along the ray
- $C'_i = \max_j(C_i, C'_{i-1})$
- $A'_i = A_j$

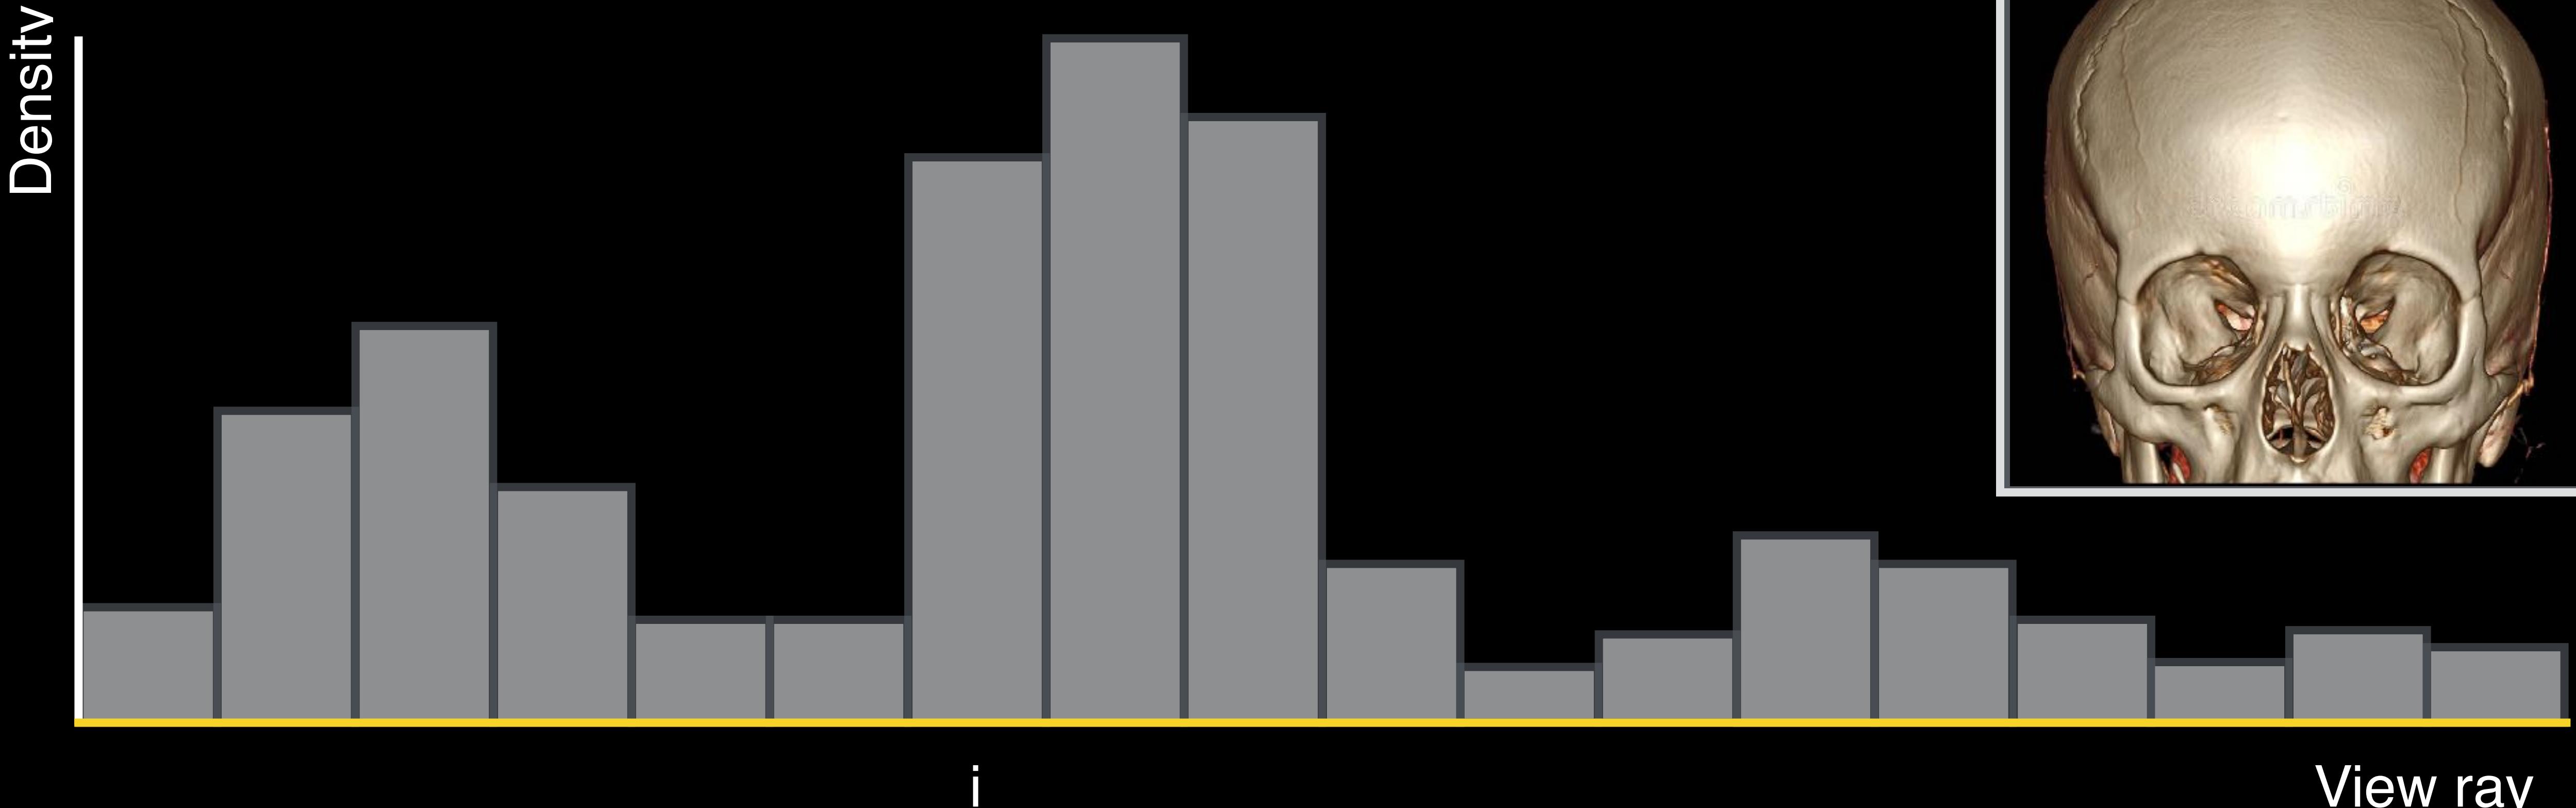
C: Color

A: Transparency/Alpha

Prime: Accumulated

Non-prime: Current value

FIRST HIT COMPOSITING



- Use the first value that is above a given threshold
- $C'_i = \text{first } C_i \text{ above threshold} := j$
- $A'_i = A_j$
- -> Isosurface rendering

C: Color

A: Transparency/Alpha

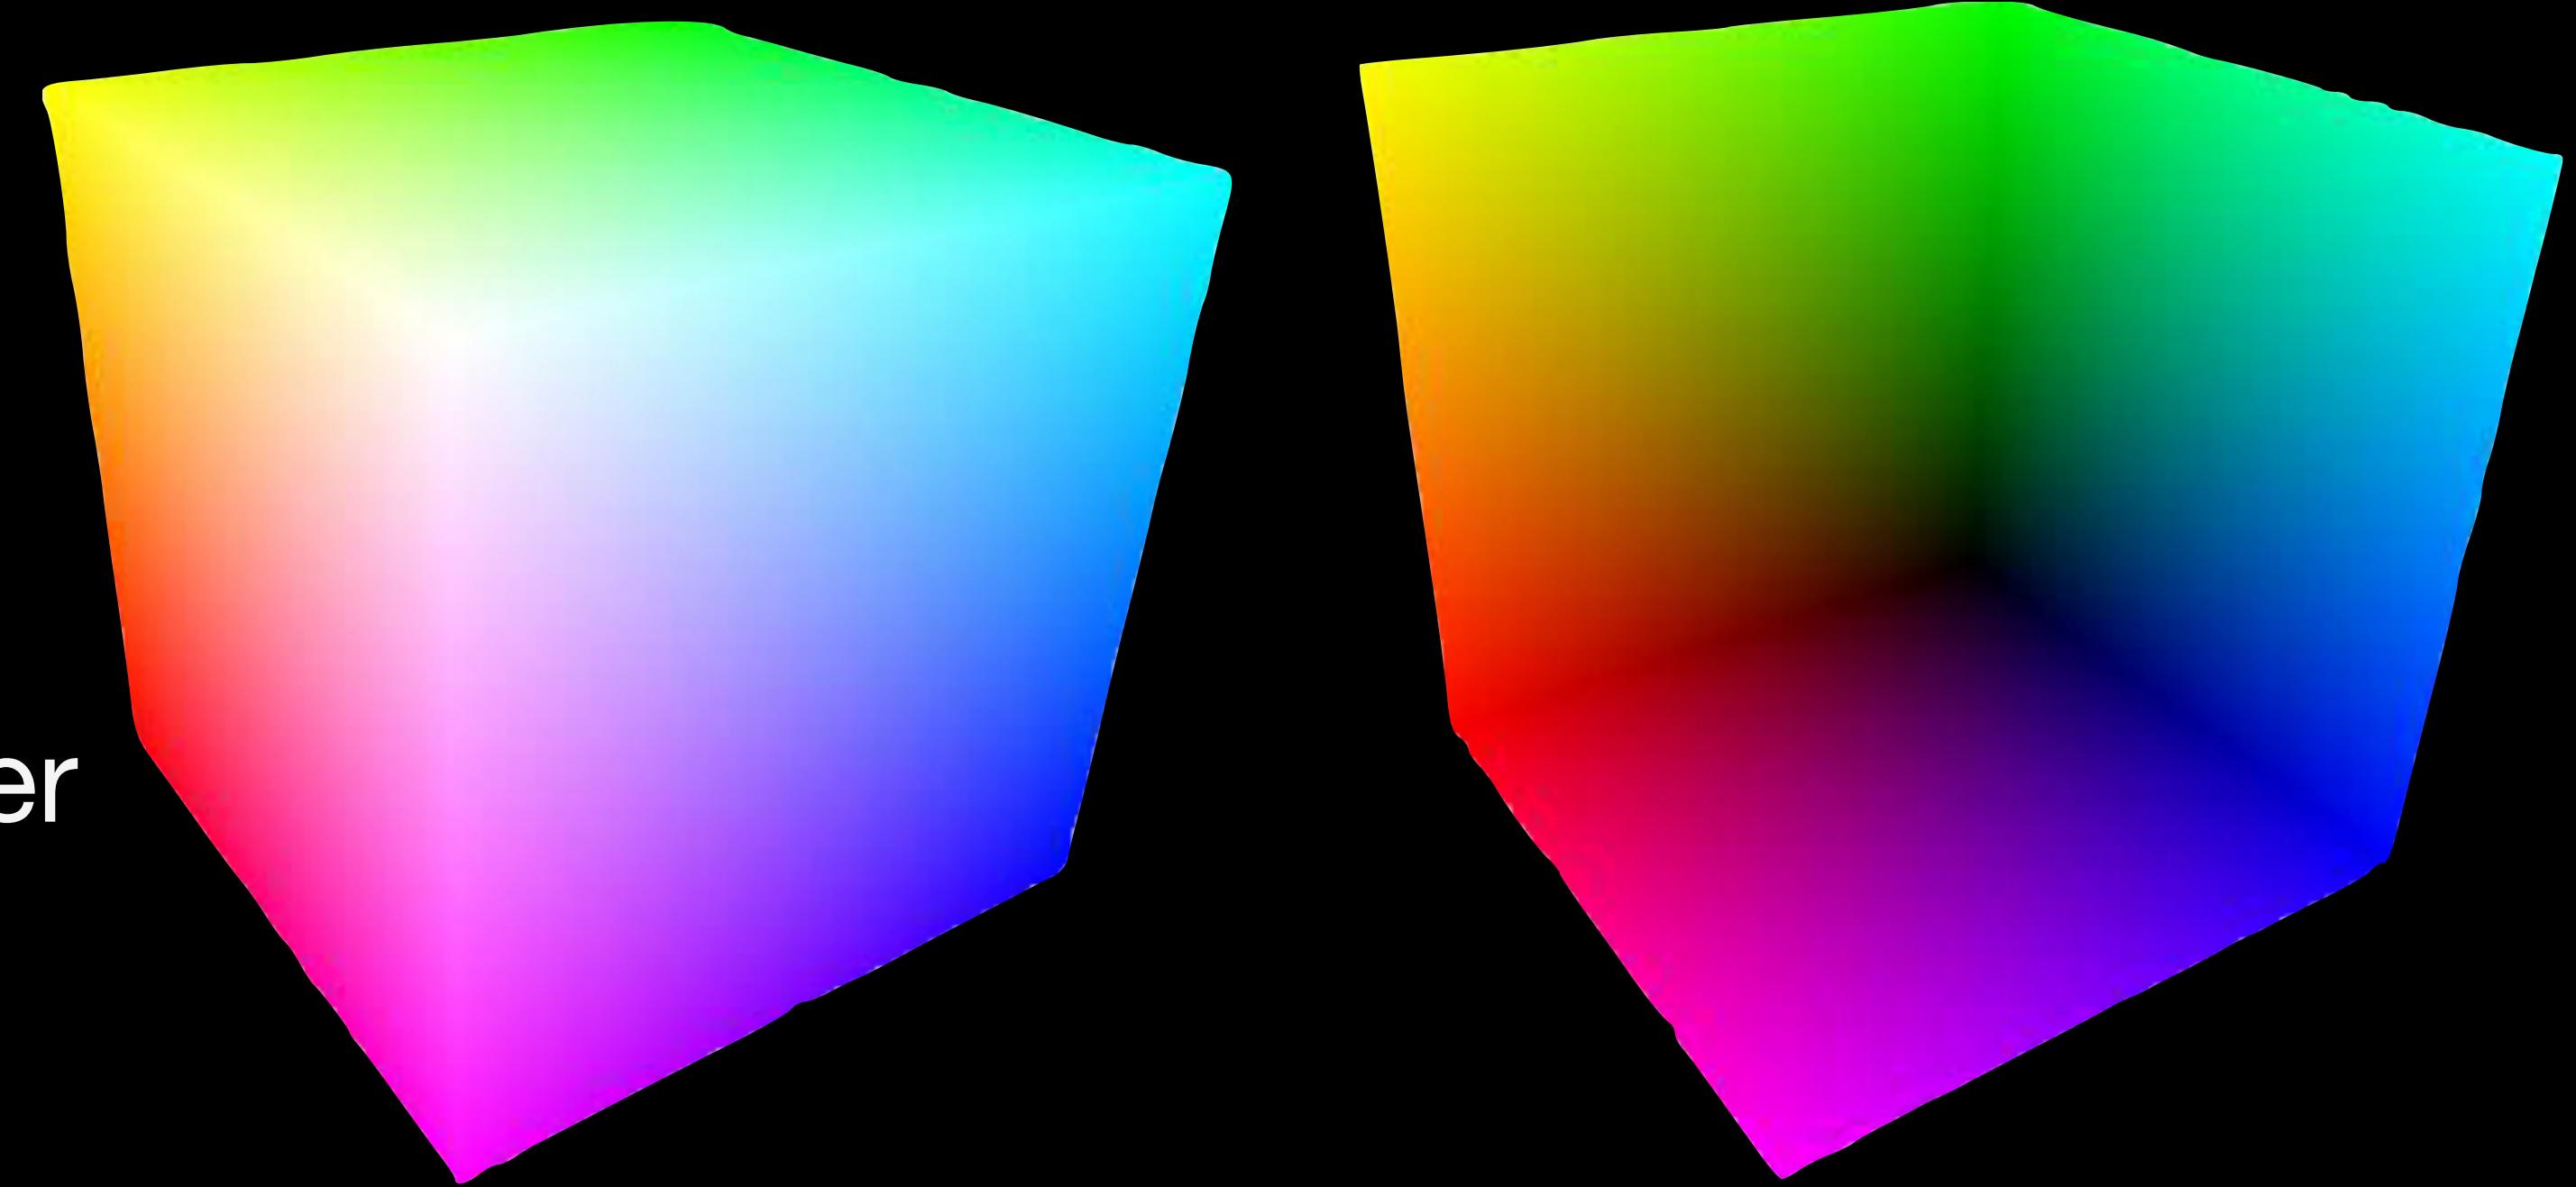
Prime: Accumulated

Non-prime: Current value

TECHNIQUES

KRÜGER-WESTERMANN PROXY GEOMETRY

- Render *Proxy Geometry* twice using the position as color [0, 1]
- First pass
 - Render front front side of the proxy geometry into a frame buffer
- Second pass
 - Render back side of the proxy geometry using the same camera parameters
- Third pass
 - Compute the difference between back side and front side color -> direction vector
 - Loop over the ray and perform compositing



Krüger, Westermann, 2003
Acceleration Techniques for GPU-based Volume Rendering

RAYMARCHING

```
vec4 traverseRay(vec3 first, vec3 last) {
    vec4 result = vec4(0.0);
    vec3 direction = normalize(last - first);
    float stepIncr = length(last - first) / numSteps;           <- uniform float numSteps
    float t = 0.0;
    for (int i = 0; i < numSteps; ++i) {
        vec3 sampleCoord = first + t * direction;
        float intensity = texture(volume, sampleCoord).a;          <- uniform sampler3D volume
        vec4 color = classify(transferFunction, intensity);          <- uniform sampler1D transferFunction
        if (color.a > 0.0) {
            result.rgb = result.rgb + (1.0 - result.a) * color.a * color.rgb;
            result.a   = result.a + (1.0 - result.a) * color.a;
        }
        t += stepIncr;
    }
    return result;
}
```

DEMO

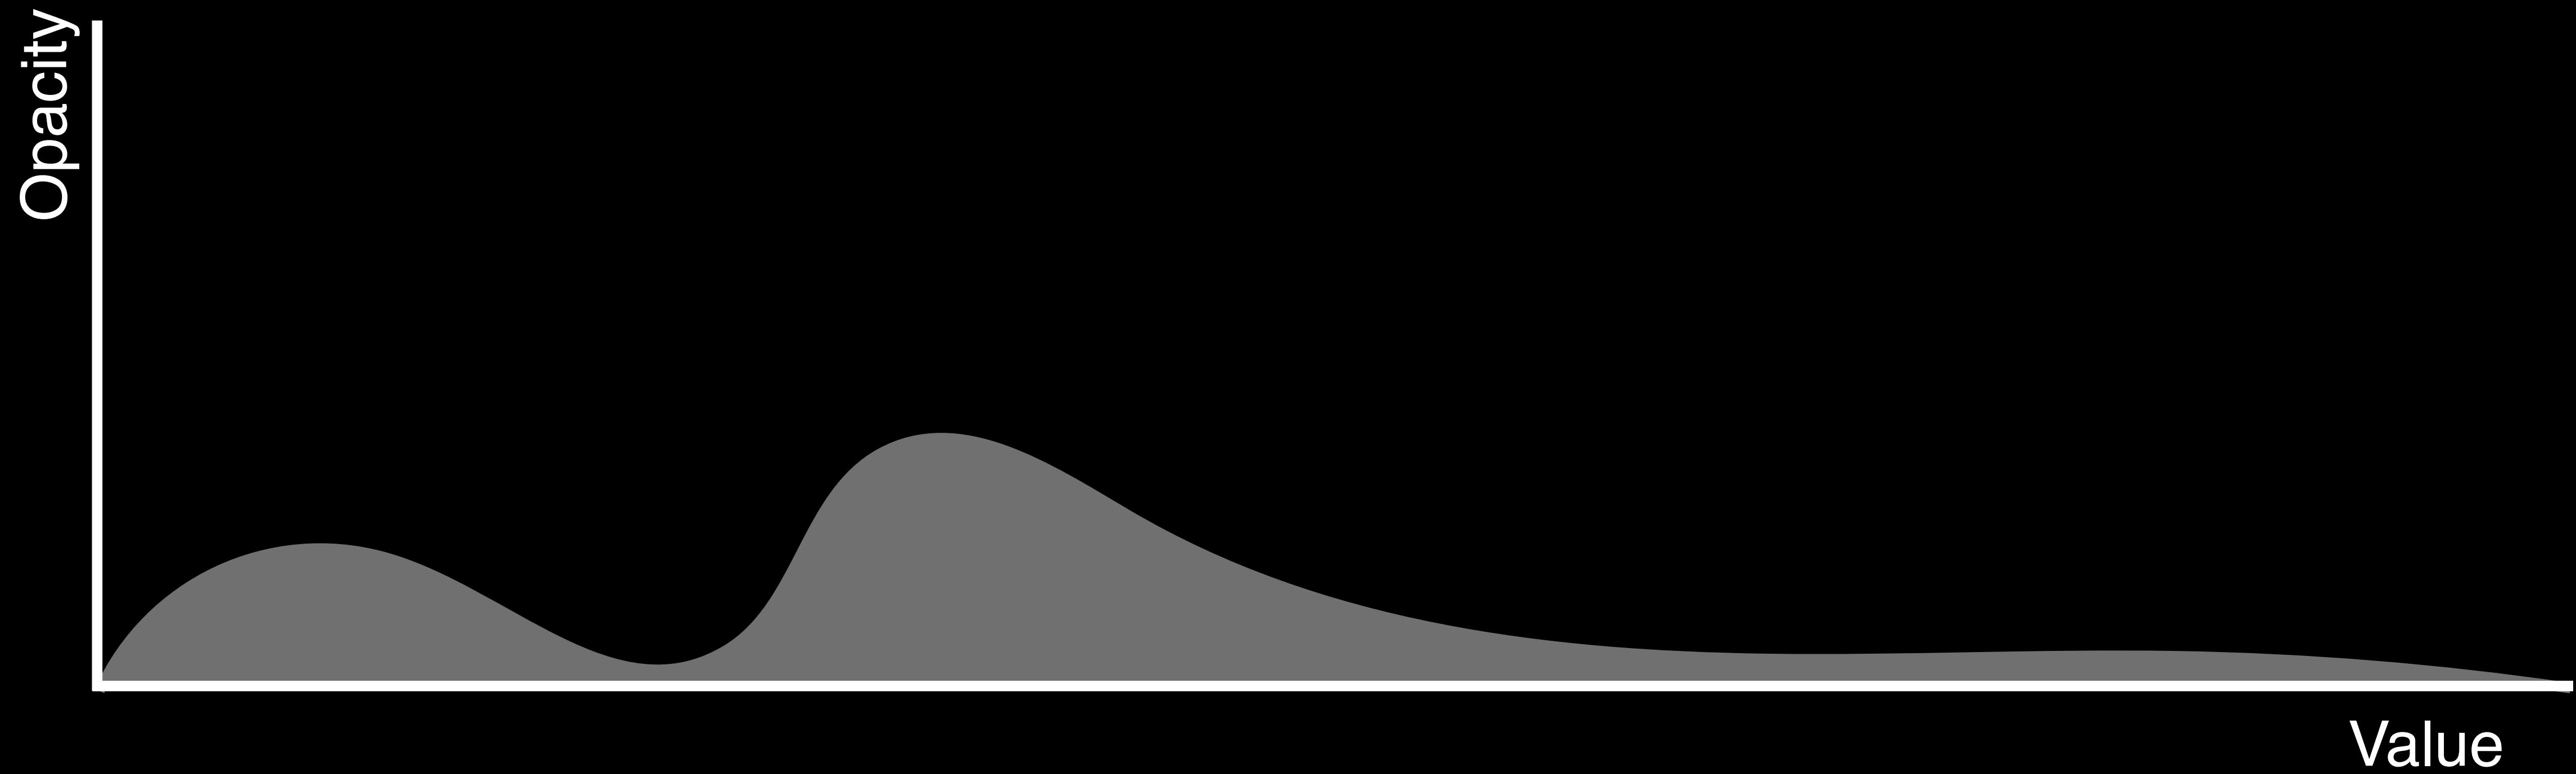
WEBGL

TRANSFER FUNCTIONS

- Most volumes have a single quantity (density, x-ray absorption, # positron emissions, ...)
- We want to create colored images
- Transfer function (or lookup table (LUT)): $f : R^i \rightarrow R^3$ (=RGB)
- Majority of cases: $i = 1$
- In OpenGL, usually implemented with a `sampler1D` that is texture sampled using the quantity as parameter (requires normalized parameters)

TRANSFER FUNCTIONS

- Many ways of specifying transfer functions



- Background in this case: Histogram over the volume

TRANSFER FUNCTIONS

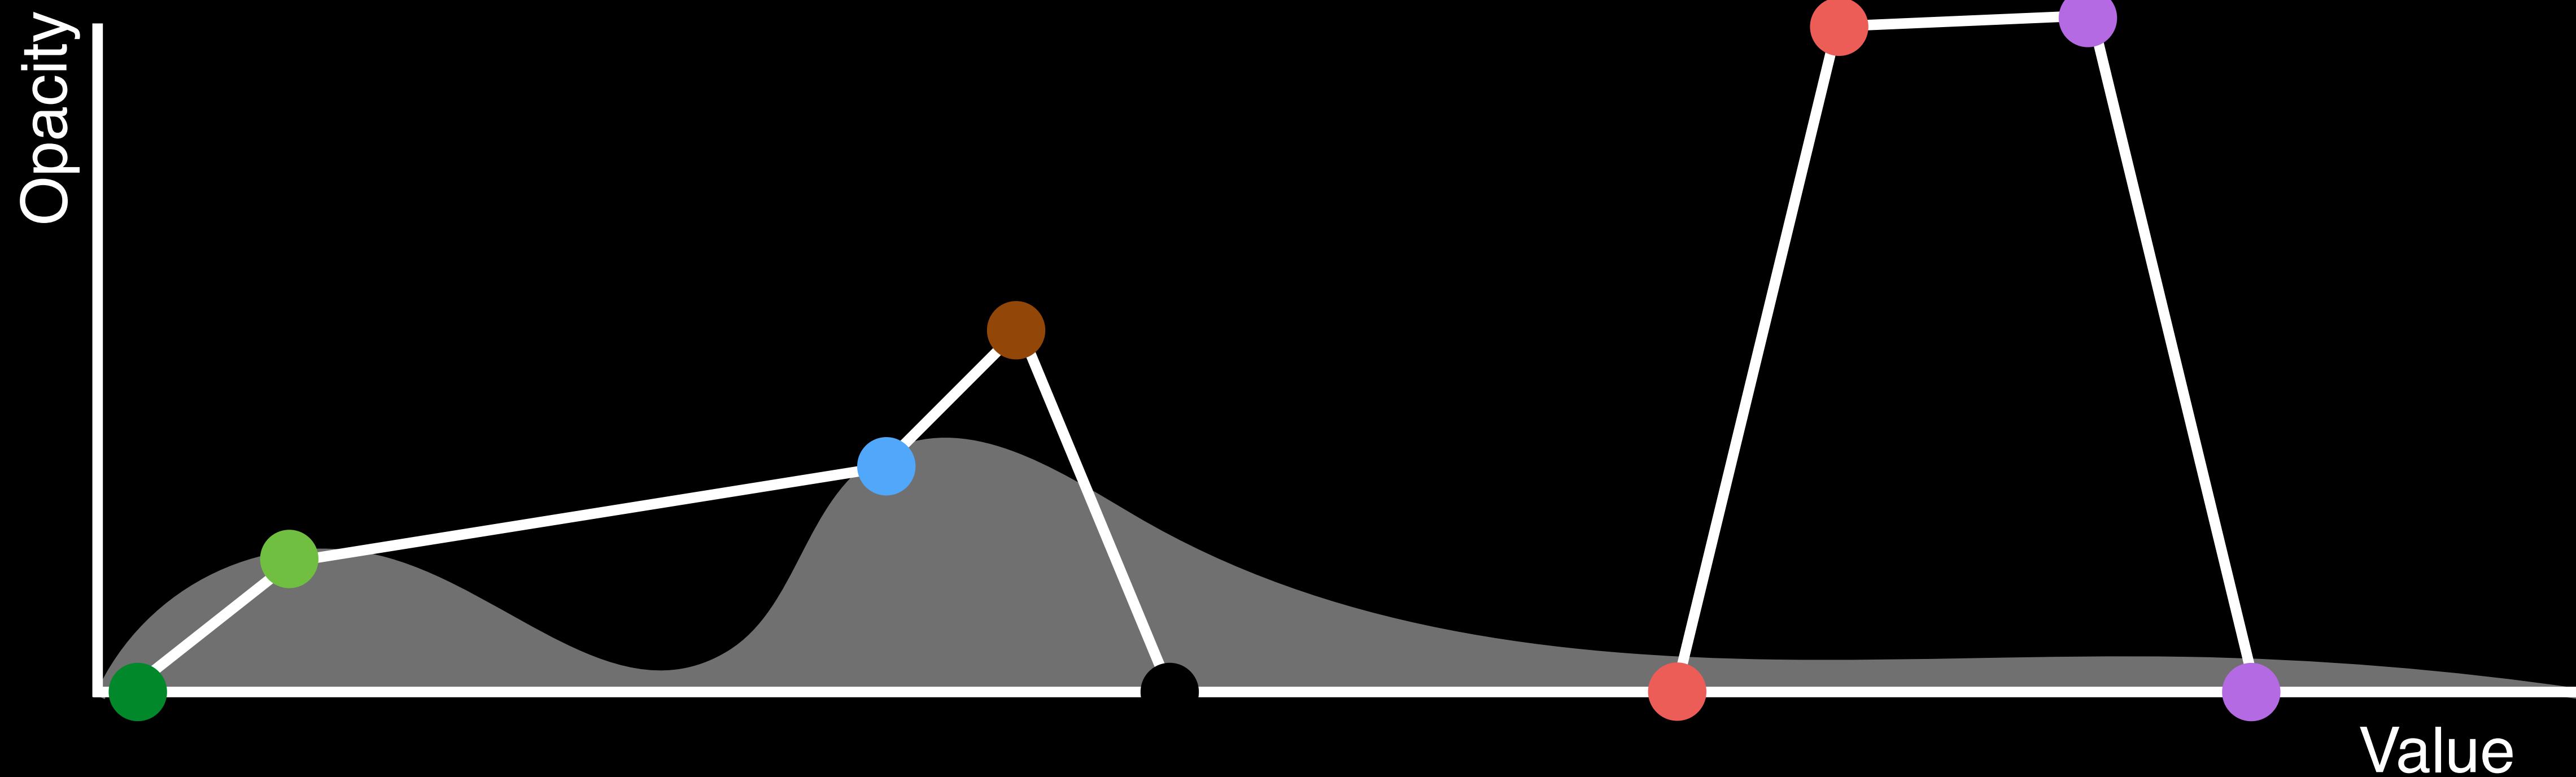
- Many ways of specifying transfer functions



- Widget based
- If a value is covered by multiple widgets, their colors are blended
- More intuitive mapping between visualized objects and colors

TRANSFER FUNCTIONS

- Many ways of specifying transfer functions



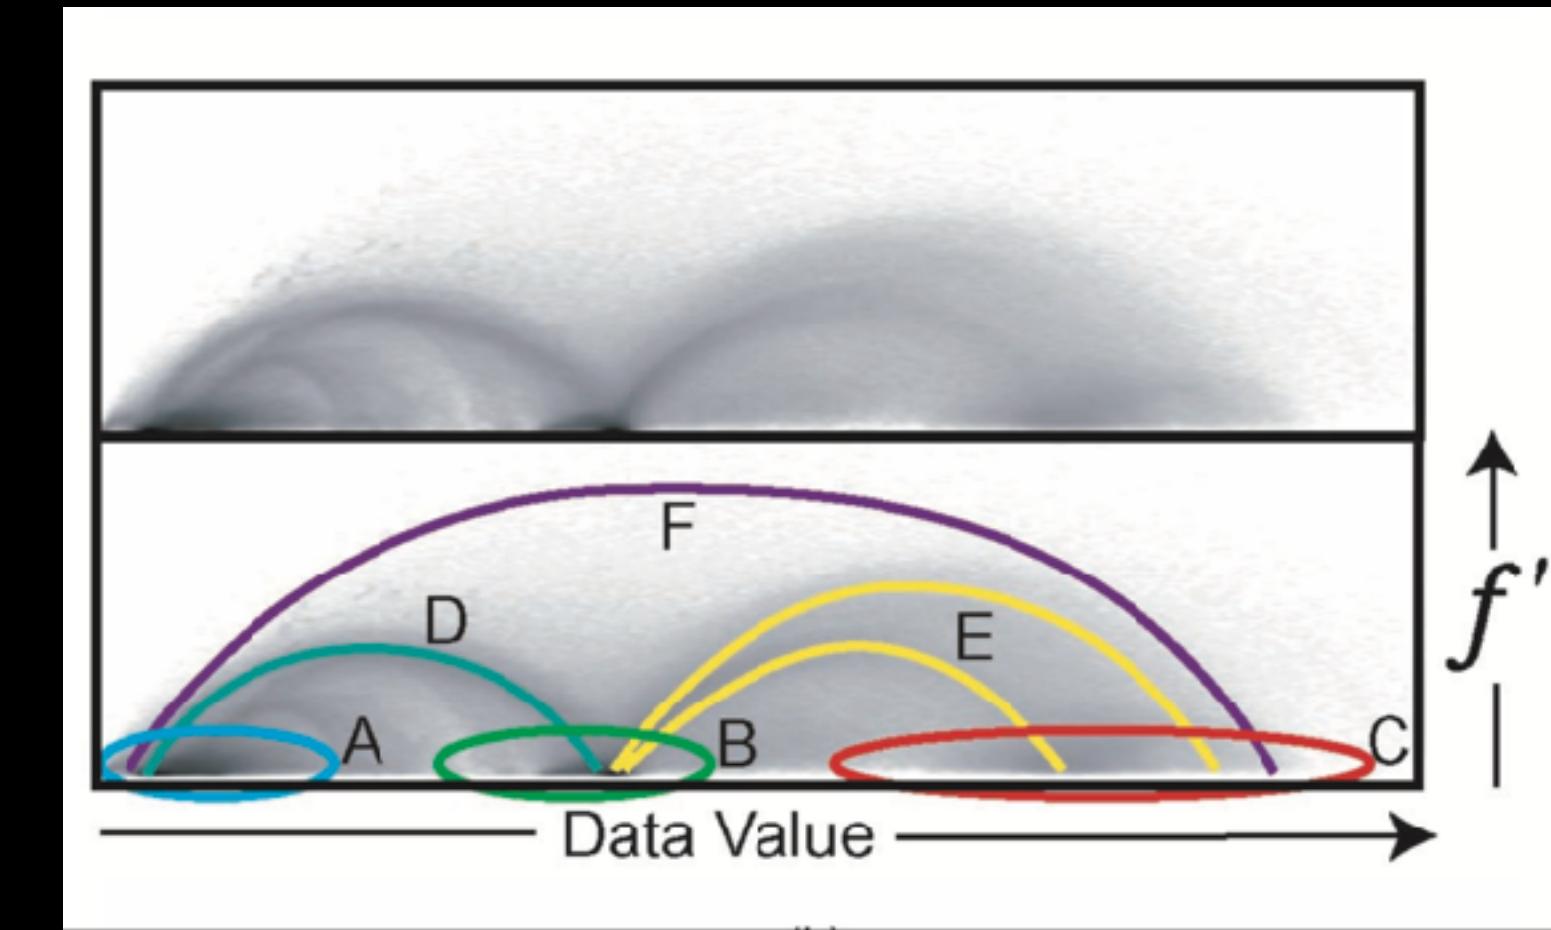
- Node based, colors interpreted between nodes
- More holistic approach

DEMO

[WWW.INVIWO.ORG](http://www.inviwo.org)

MULTIDIMENSIONAL TRANSFER FUNCTIONS

- Not often used
- Examples
 - f vs $\| \nabla f \|$
 - f vs $\| \nabla f \|$ vs $(1 / \| \nabla f \|)^2 * (\nabla f)^T * Hf * \nabla f$
 - Principle curvature



Kniss, Kindlmann, Hansen, 2002
Multidimensional Transfer Functions
for Interactive Volume Rendering

LIGHTING MODELS

- Local lighting models similar to OpenGL lecture (Blinn-Phong)
- Problem: Since it is **direct** volume rendering, no explicit surface exists, therefore no normal
- -> replacement: gradient ∇f as replacement for the normal in lighting calculations
- Approximations of ∇f using finite differences
 - Forward difference, backward difference, *central difference*
 - Central difference:
 - $f'(x,y,z) = f(x+h, y, z) - f(x-h, y, z),$
 $f(x, y+h, z) - f(x, y-h, z),$
 $f(x, y, z+h) - f(x, y, z-h)$

GLOBAL ILLUMINATION

- Local lighting models cannot represent shadows, caustics, refractions, and other global effects
- *Global Illumination* techniques such as *Photon Mapping*, *Metropolis light transport*, *Radiosity*, and *raytracing* provide higher quality shadows, but are slower to compute



Image © Daniel Jönsson

GLOBAL ILLUMINATION

- For each sample evaluate the occlusion towards the light source
- In the case of *raytracing*, this is done through secondary rays that sample the volume analogous to the primary ray

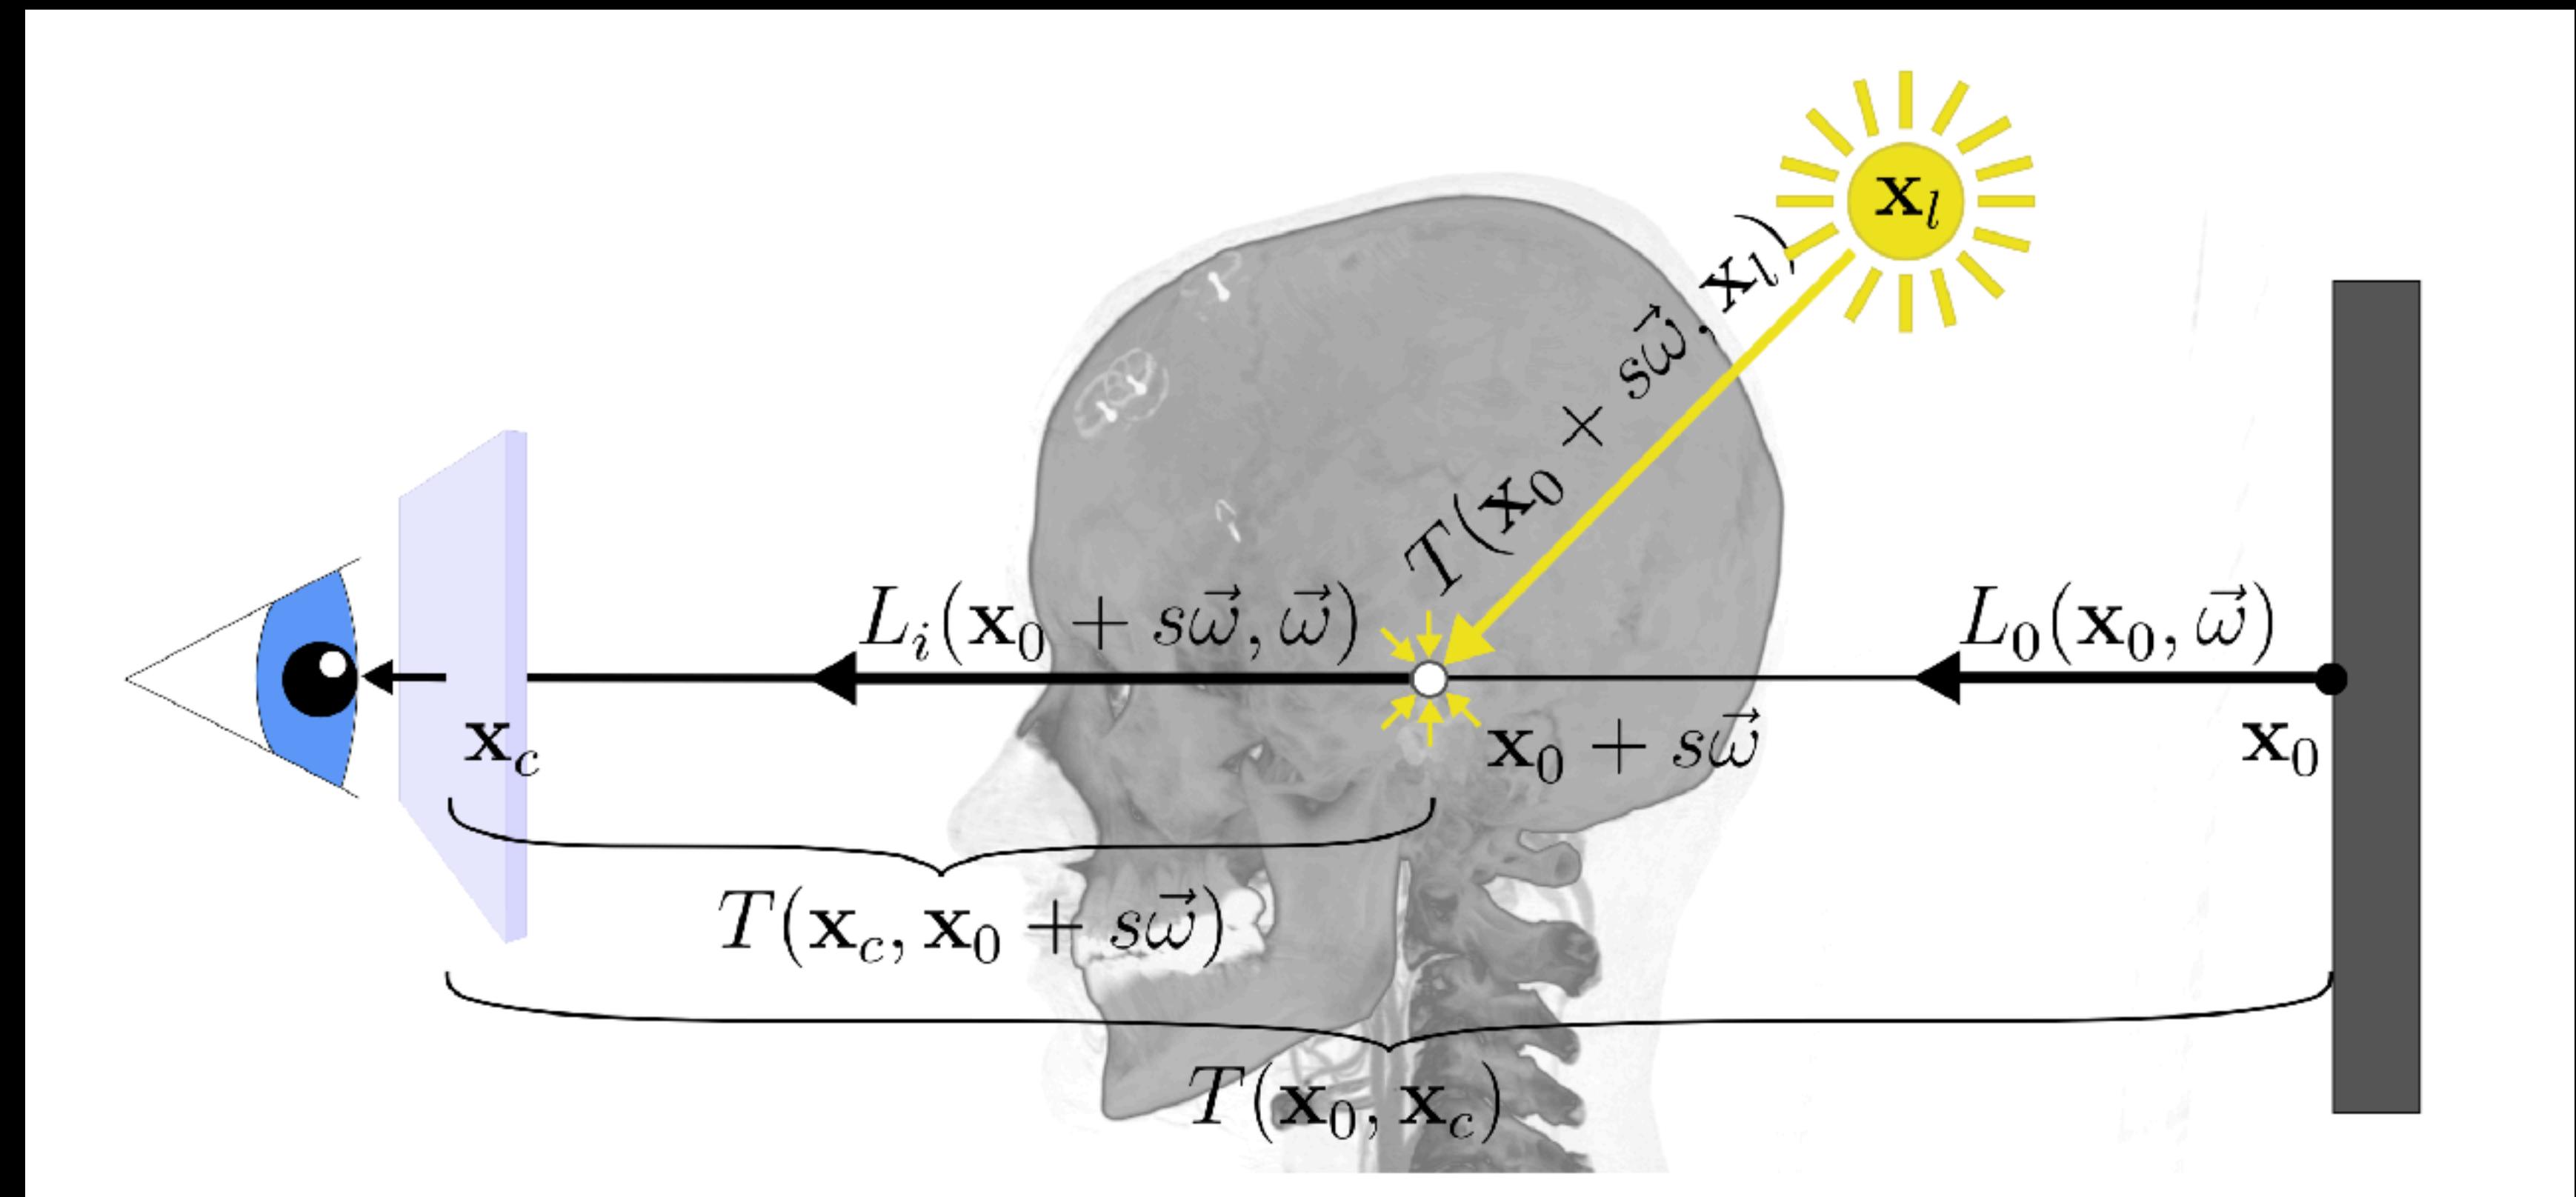
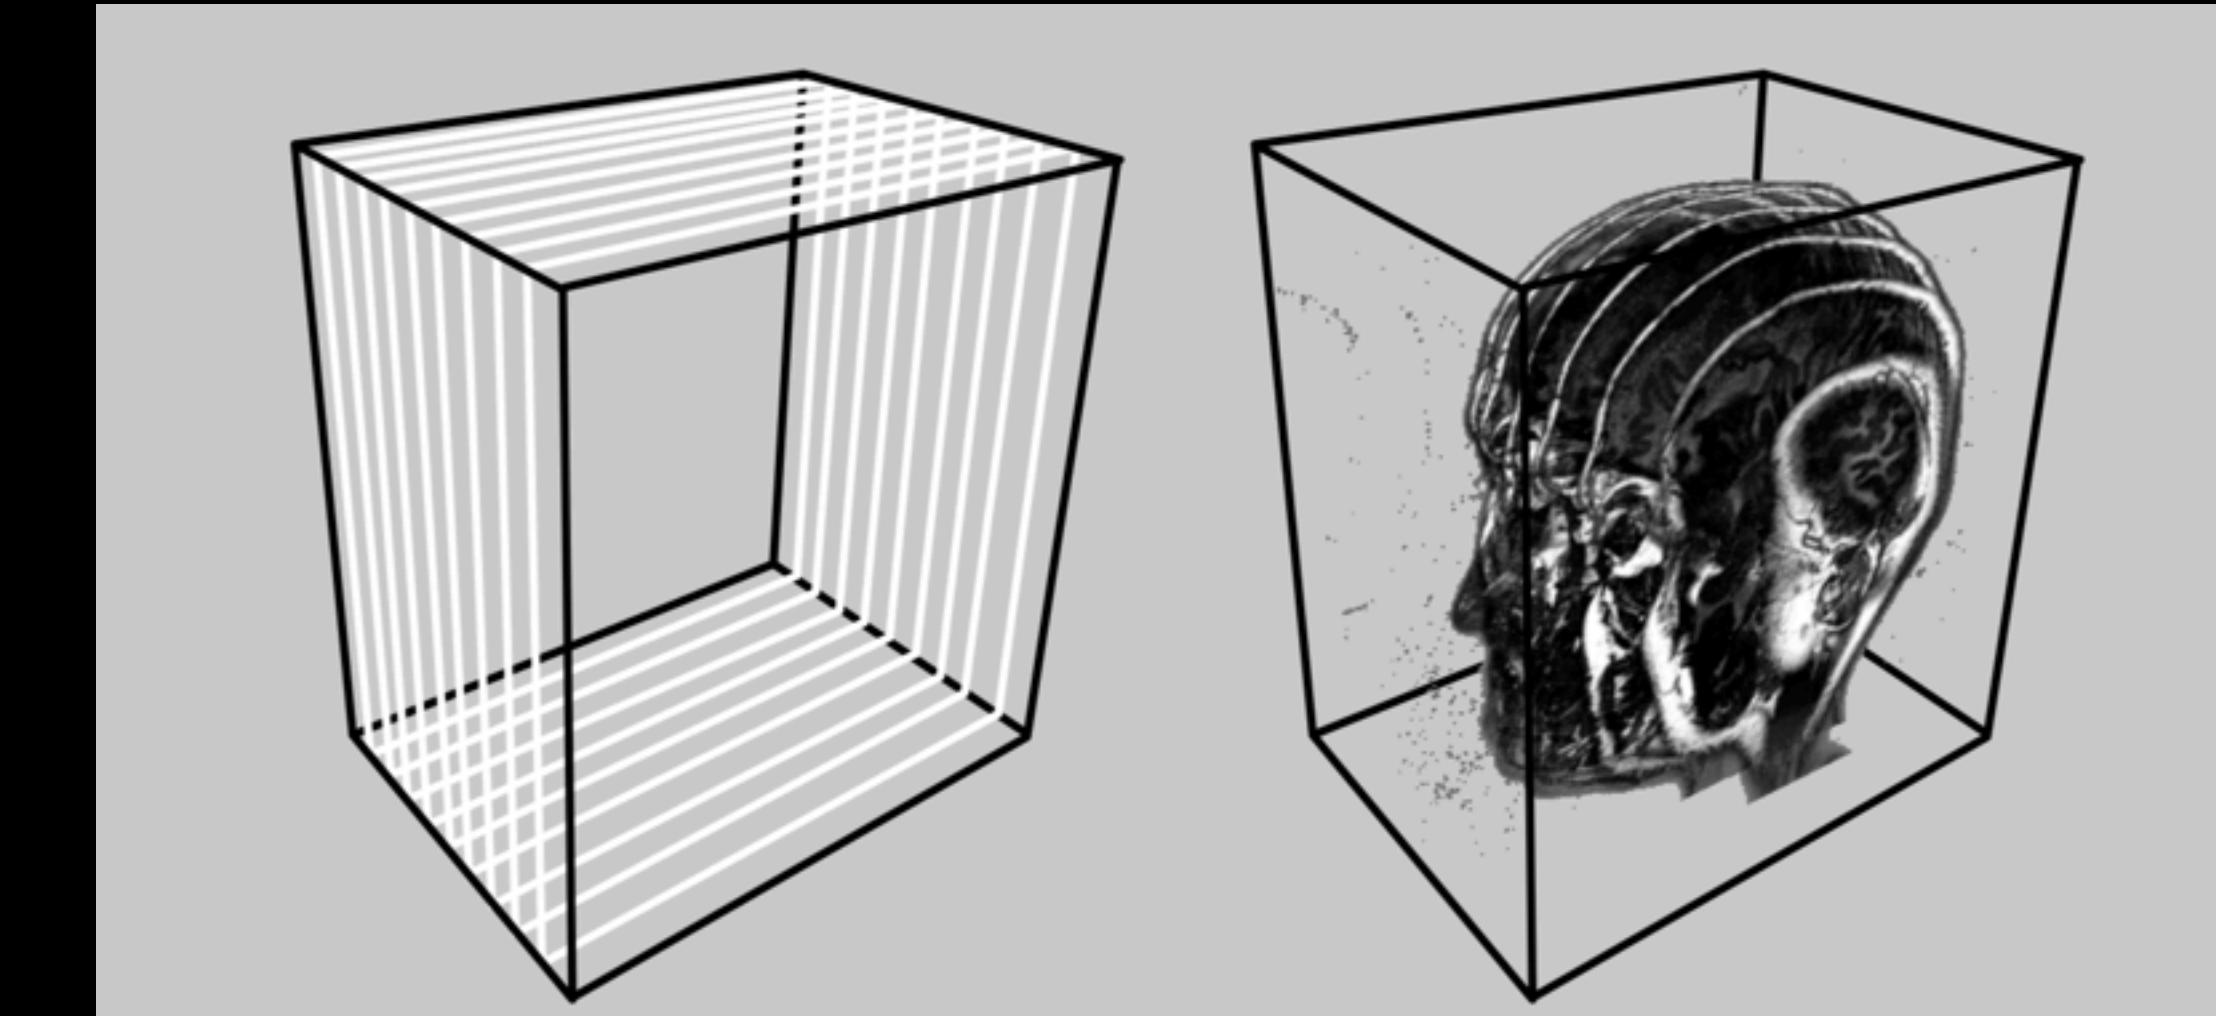


Image © Daniel Jönsson

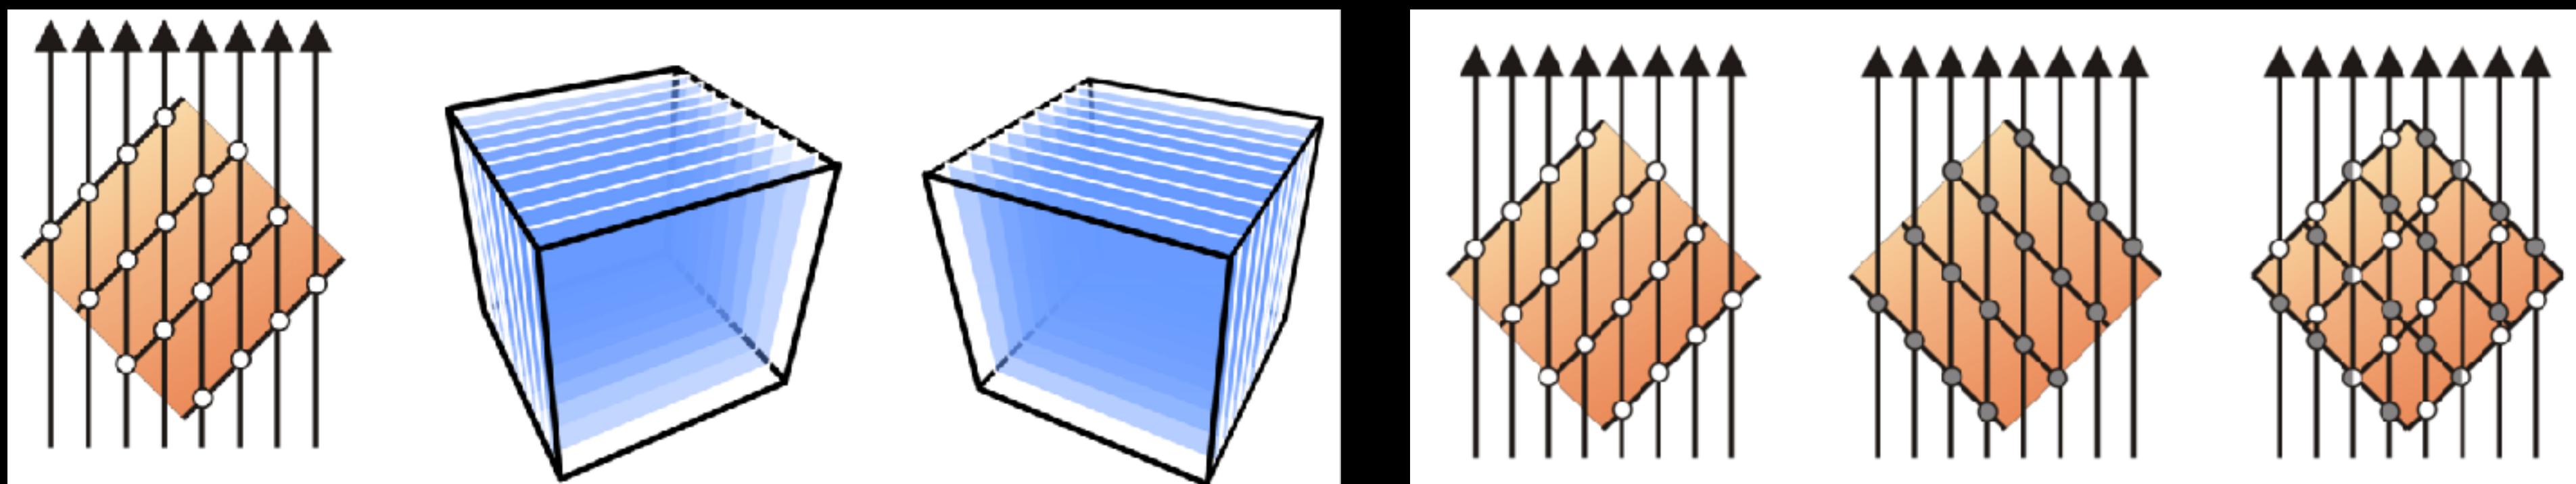
OBJECT-ORDER METHODS

TEXTURE SLICING

- Representing volume as stacks of 2D textures that are composited



- Option 1: Three stacks of 2D textures along primary axes
- Problematic if view direction changes past 45° as sampling points change drastically
- Does not require any complex shader programs; no need for 3D texture support

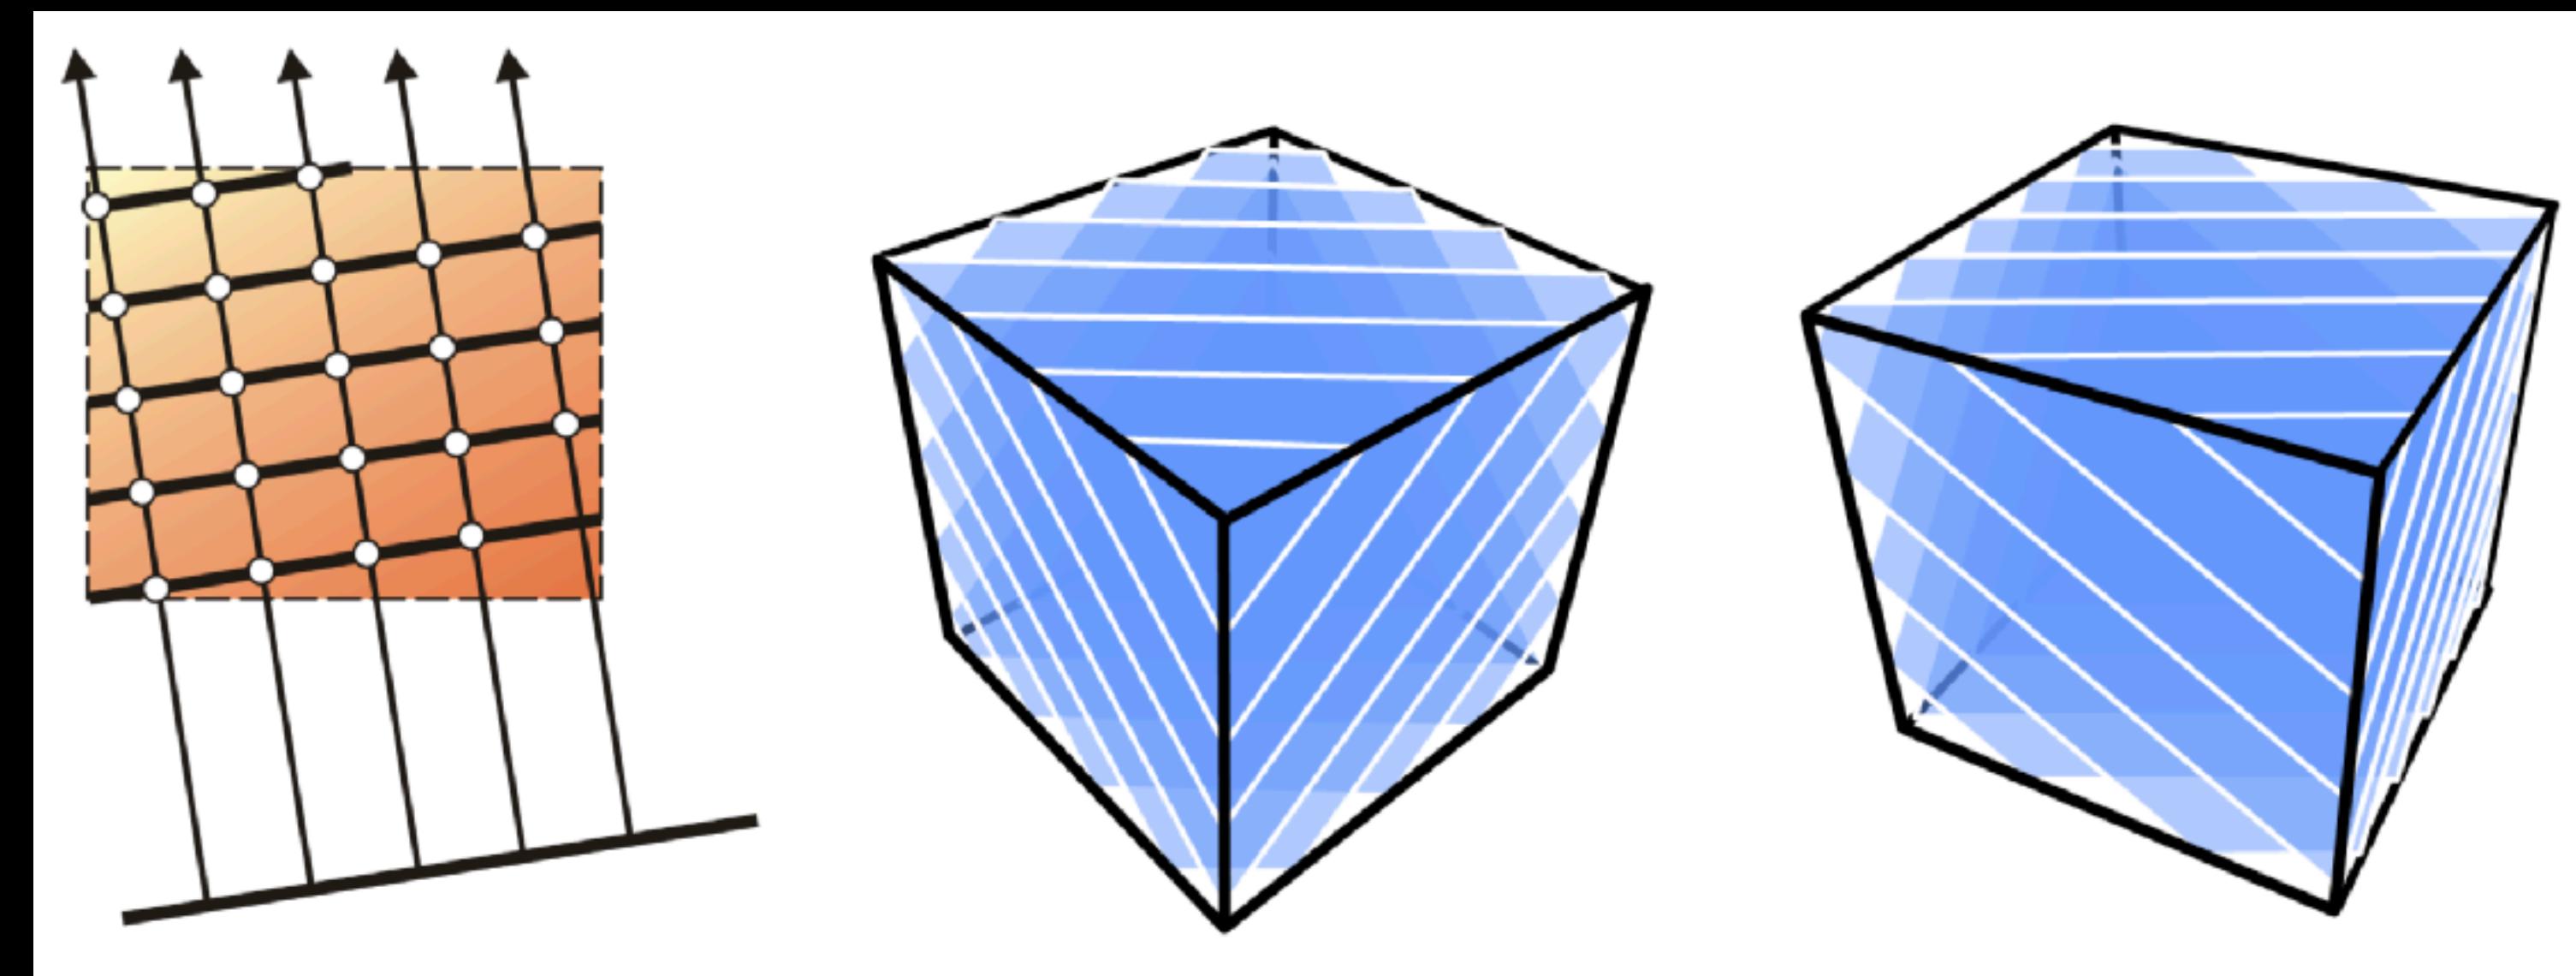


Images © Weiskopf, Machiraju, Möller

TEXTURE SLICING

- Option 2: Single stack of 2D textures along view direction

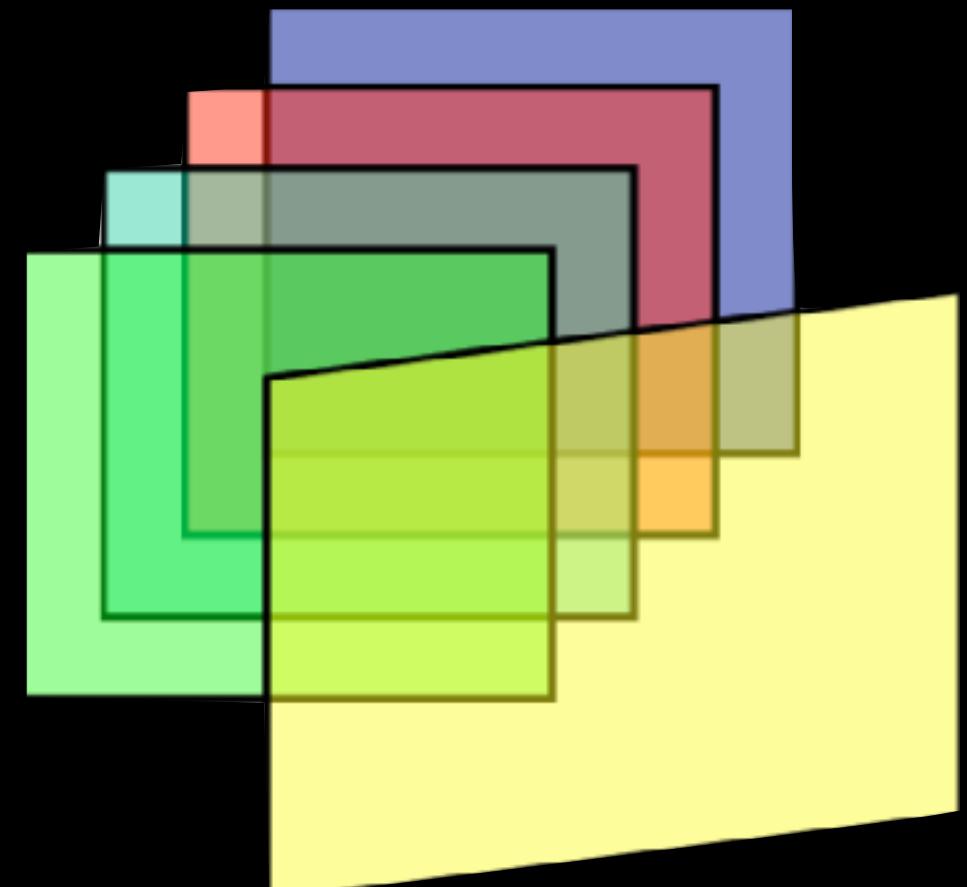
- Trilinear interpolation slower than bilinear interpolation
- Sampling distance independent of viewing angle



Images © Weiskopf, Machiraju, Möller

VOLUME SPLATTING

- Object-order
- Project each voxel onto the image plane individually
- Each voxel is represented by a 3D kernel
 - Kernel is converted into 2D footprint on the image plane
 - Size and shape of kernel determines image quality (sharpness, # holes, ...)
- One voxel is splatted onto many pixels
- Voxels are added within sheets
 - Front-to-back compositing each sheet
 - Accumulating final result in a separate buffer



Westover, 1990

Splatting: A Parallel, Feed-Forward Volume Rendering Algorithm